



ACID in SQL Server

what it is, how it works, and how to live more adventurously.

In this session

- What is ACID compliance
- Isolation levels in SQL Server
- Durability in SQL Server

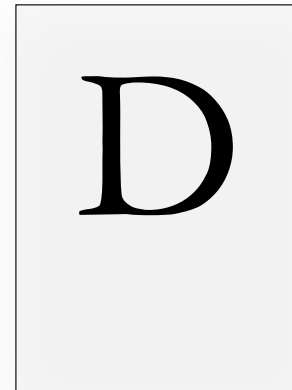
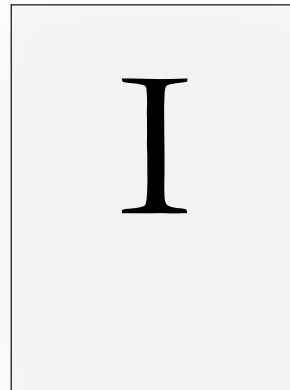
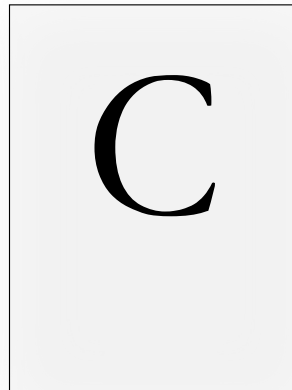
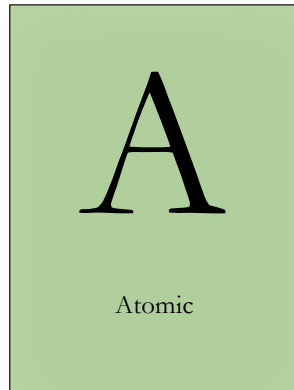


A





C

I





D





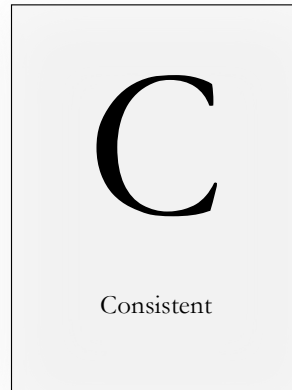
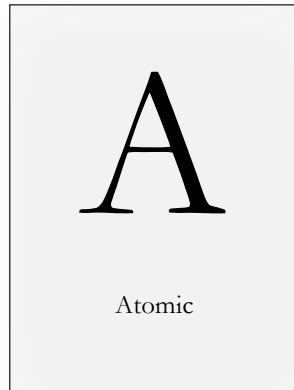
All or nothing.
The *business* logic.



A single truth.
The *database* logic.



No interference
when changing data.



Keep all promises.

Who am I?

- SQL Server developer since 1997
- Consultant for 25+ years
- Organizer of Data Saturday Stockholm
- Data Platform MVP
- Dog person

@dhmacher on (almost) all the socials.




Isolation: Quirks and features of multi-user databases



Dirty reads

- no Nobel literature prize for you.

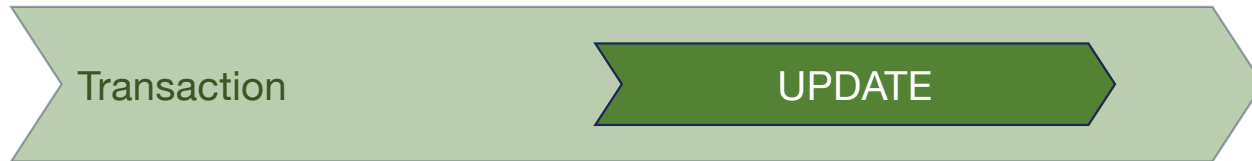


Read Uncommitted (a.k.a. NOLOCK)



Dirty reads

- no Nobel literature prize for you.



Read Committed




Dirty reads

- no Nobel literature prize for you.

- DML does not respect READ UNCOMMITTED, and cannot make changes to uncommitted data.

Dirty reads

- no Nobel literature prize for you.



Read Uncommitted (a.k.a. NOLOCK)




DEMO



Non-repeatable reads



Non-repeatable reads



DEMO



Phantom reads

“Where did that come from?”

Phantom reads

“Where did that come from?”



Phantom reads

“Where did that come from?”

Serializable



DEMO



What are isolation levels?



What are isolation levels?


Read uncommitted

Read committed

Repeatable read

Serializable

← less likely to block,
more concurrent



more likely to block, less concurrent →

← less likely to block, more concurrent

more likely to block, less concurrent →

What are isolation levels?

Read uncommitted

Read committed

Repeatable read

Serializable

Dirty reads

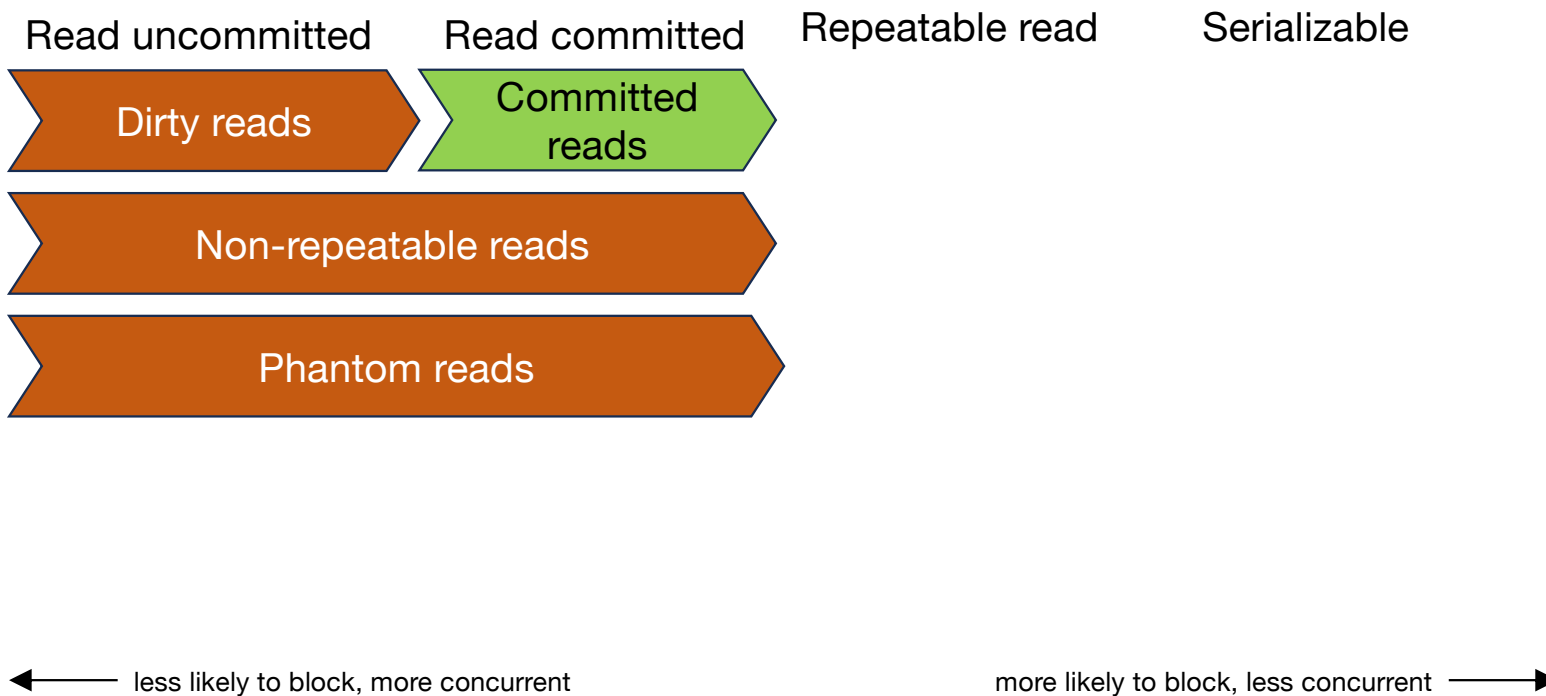
Non-repeatable
reads

Phantom reads

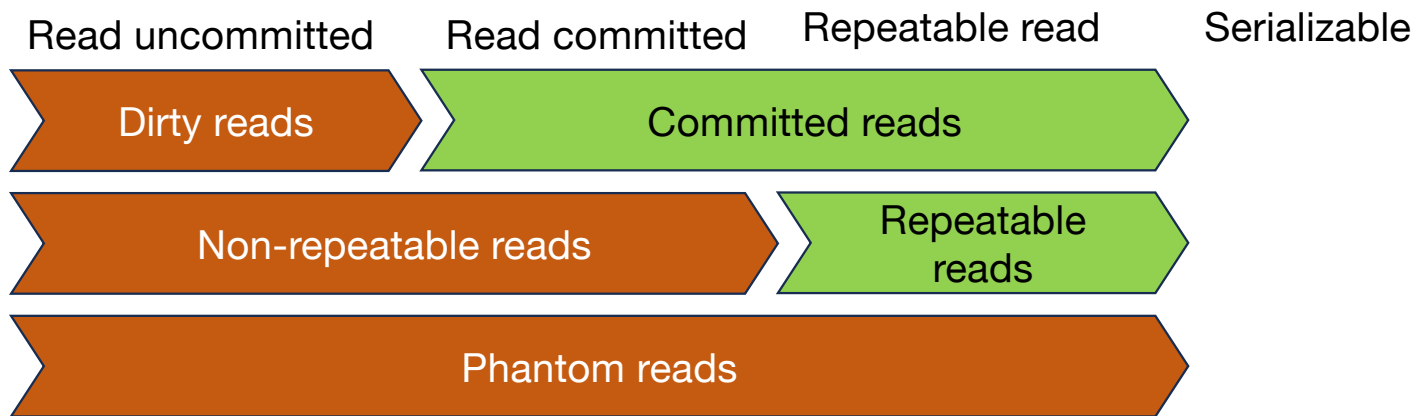
← less likely to block, more concurrent

more likely to block, less concurrent →

What are isolation levels?



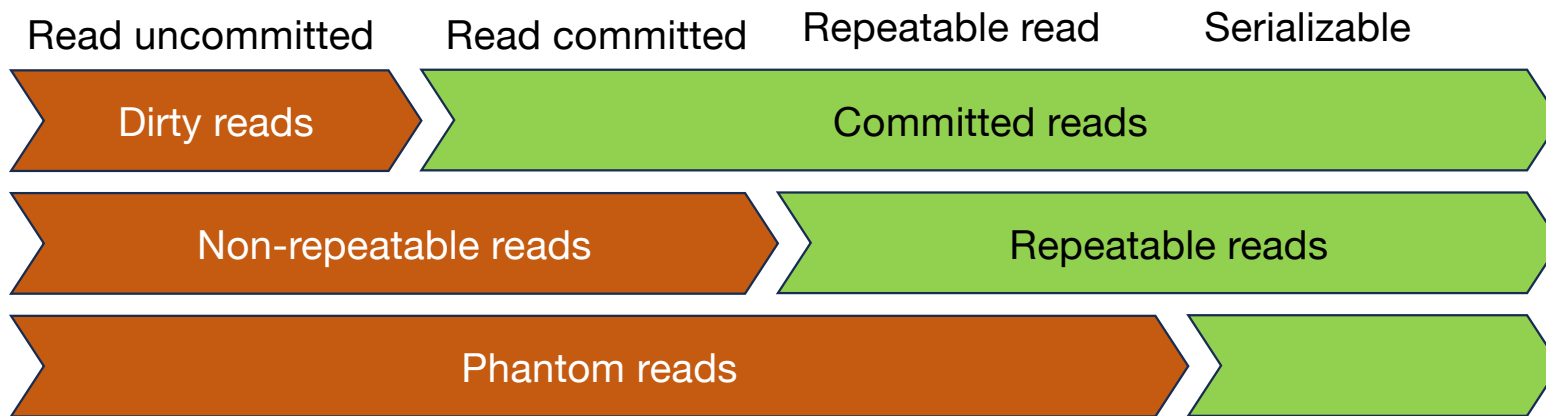
What are isolation levels?



← less likely to block, more concurrent

more likely to block, less concurrent →

What are isolation levels?




← less likely to block, more concurrent

more likely to block, less concurrent →

So... Serialize all the things, then?

Less locking

- Less isolation
- Less predictable




More locking

- Better isolation
- More predictable

So... Serialize all the things, then?

Less locking



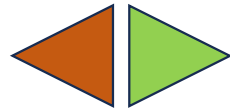
More locking

- Less isolation
- Less predictable
- Better concurrency

- Better isolation
- More predictable
- Lower concurrency

So... Serialize all the things, then?

Less locking

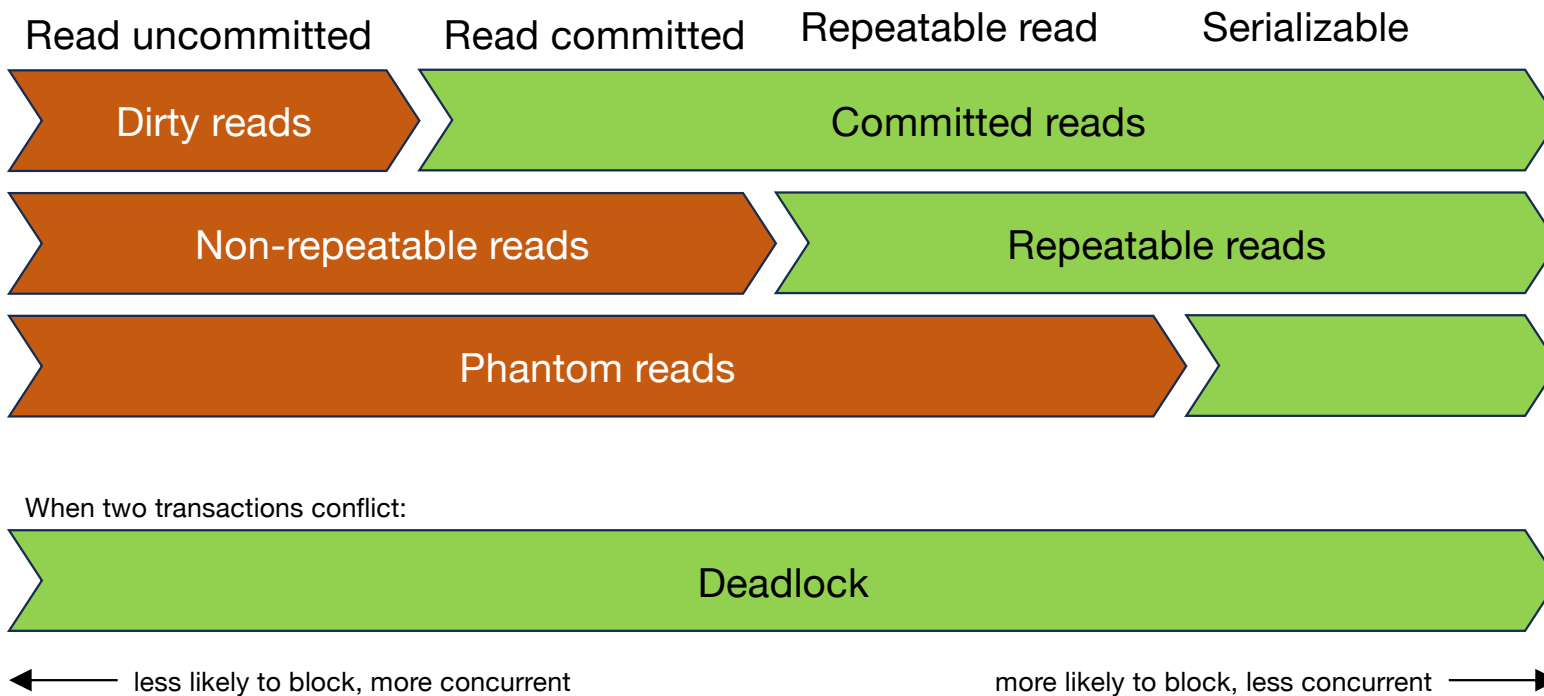


More locking

- Less isolation
- Less predictable
- Better concurrency
- Fewer conflicts

- Better isolation
- More predictable
- Lower concurrency
- More conflicts

What are isolation levels?



Deadlock



Deadlock

Transaction 1

UPDATE a

Transaction 2

Deadlock

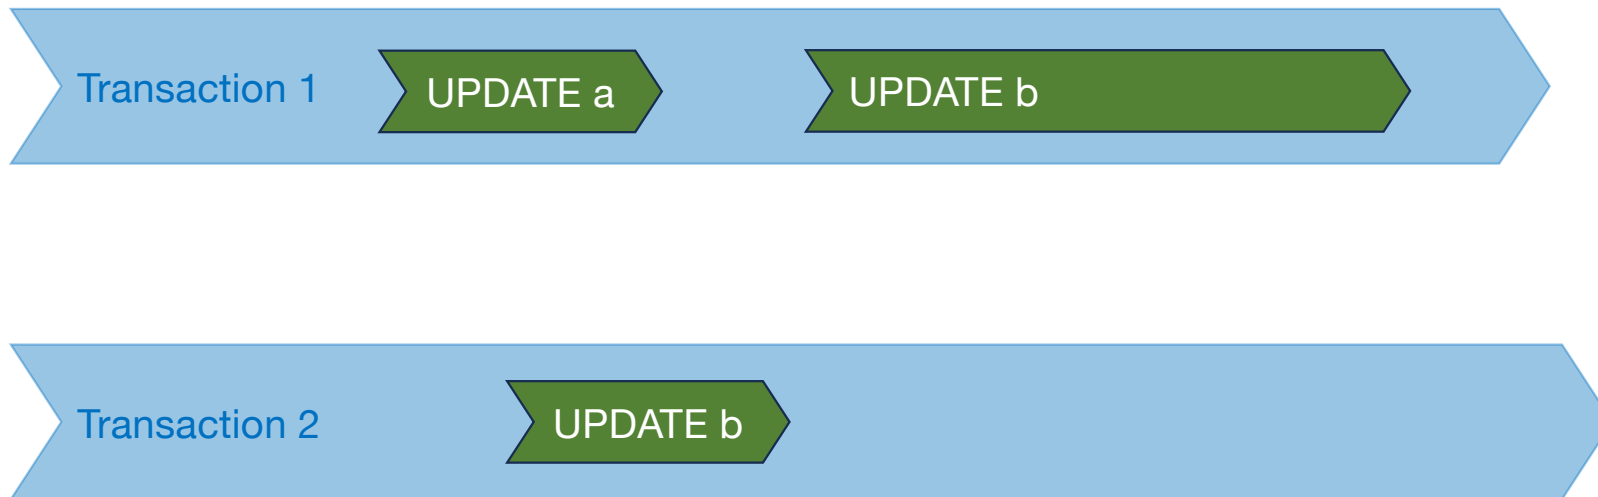
Transaction 1

UPDATE a

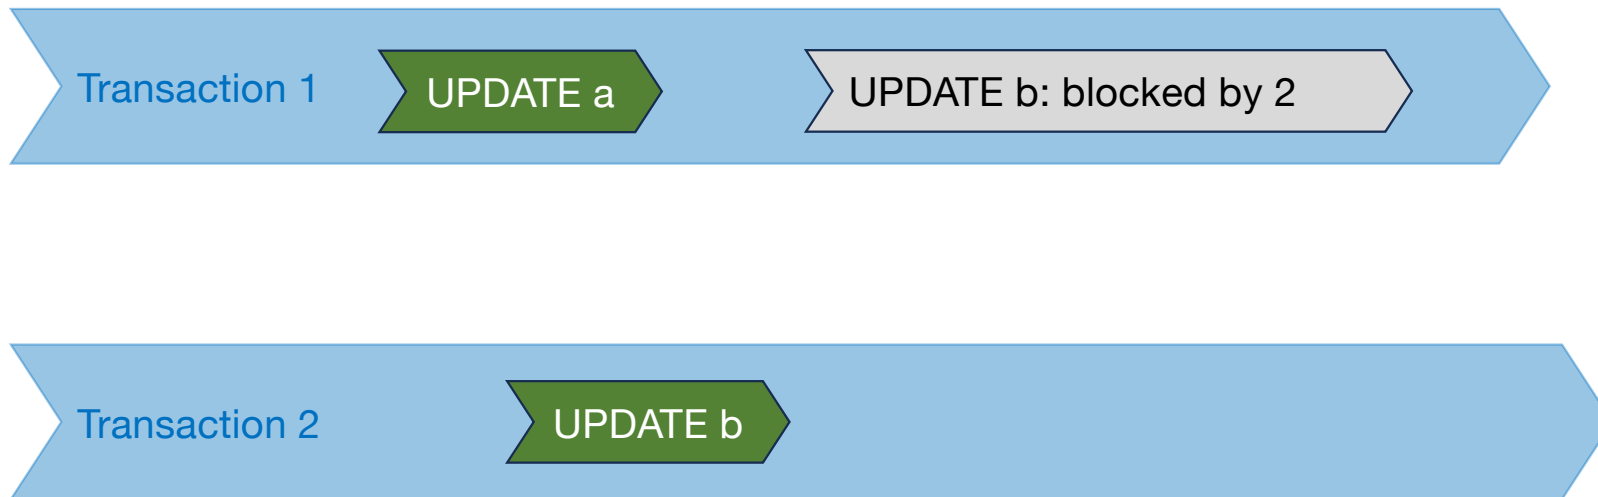
Transaction 2

UPDATE b

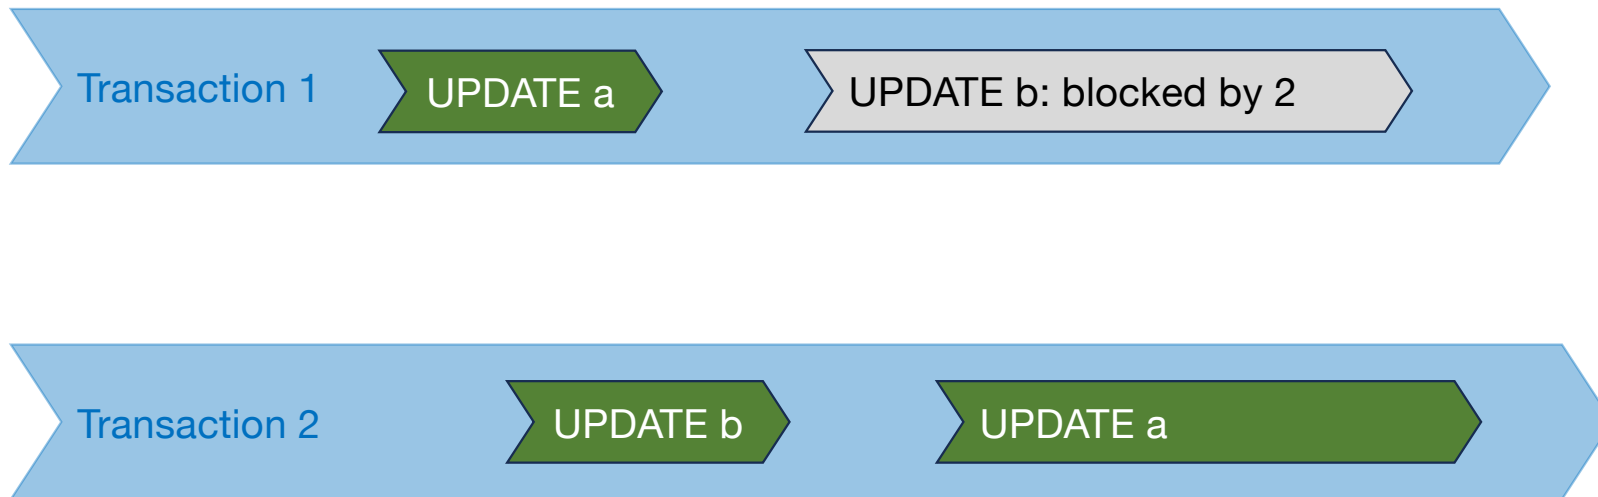
Deadlock



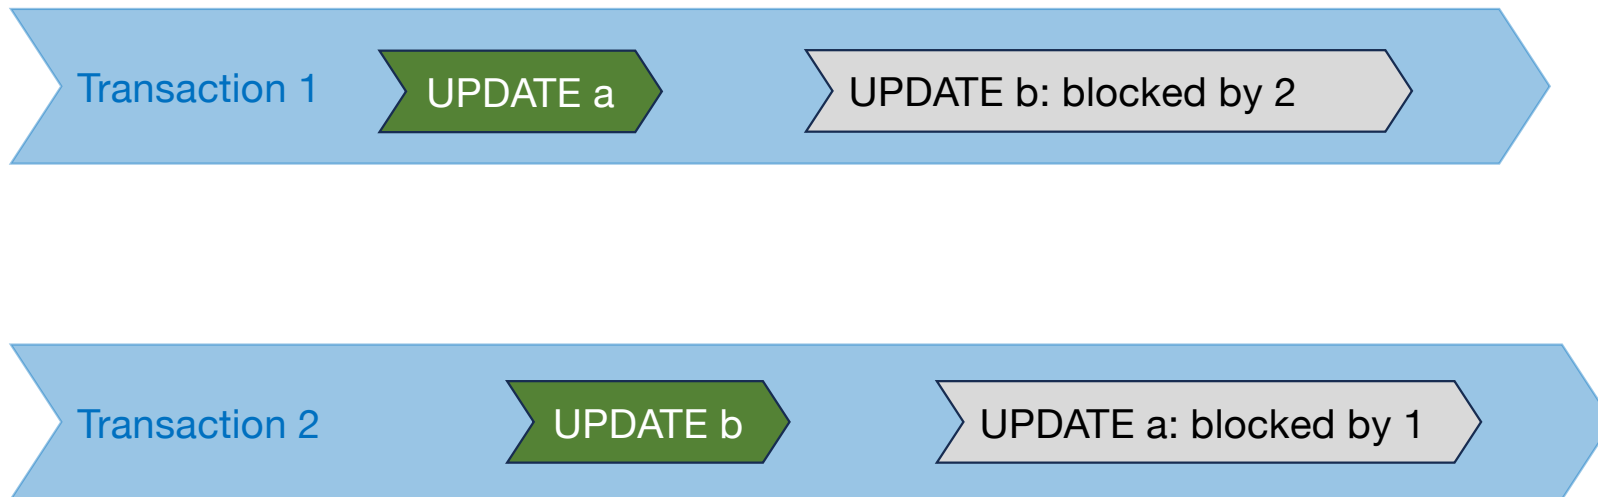
Deadlock



Deadlock



Deadlock



Deadlock

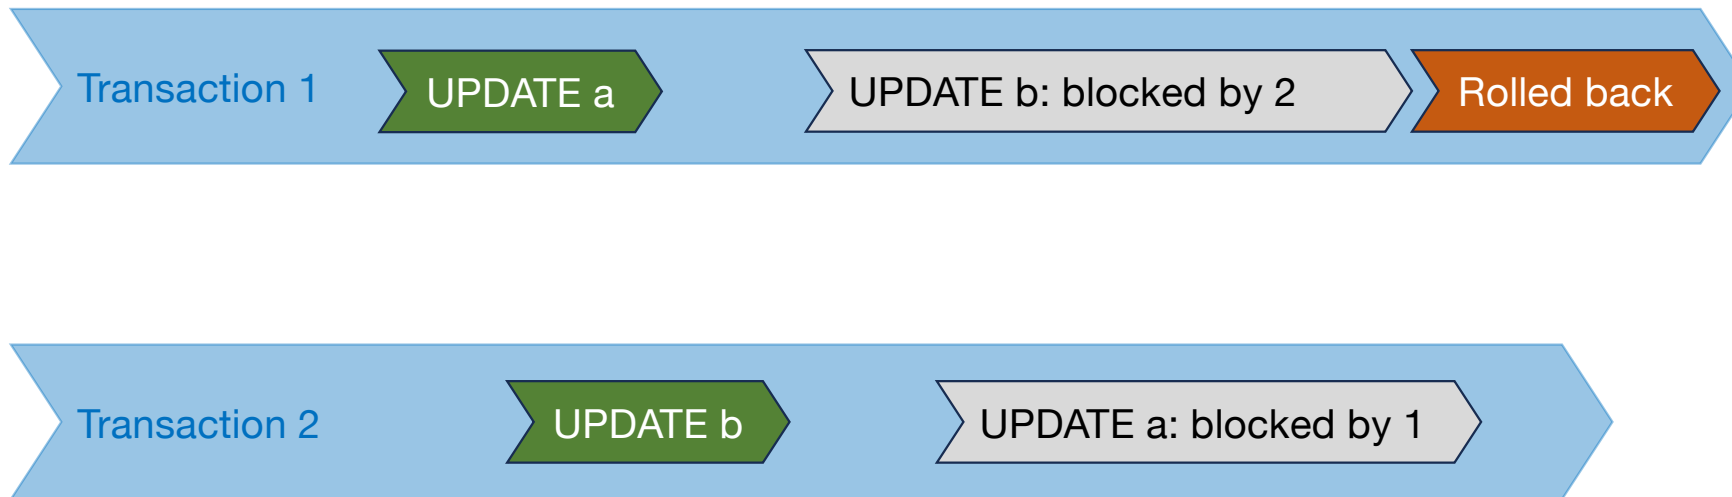


Deadlock: Transaction 1 and 2 are waiting on each other.



Deadlock

SQL Server will fail the transaction with the least amount of work to roll back:



Deadlock

Transaction 1

Rolled back

Transaction 1 is no longer blocking a.

Transaction 2

UPDATE b

UPDATE a

Deadlock

Transaction 1

Rolled back

Transaction 2

UPDATE b

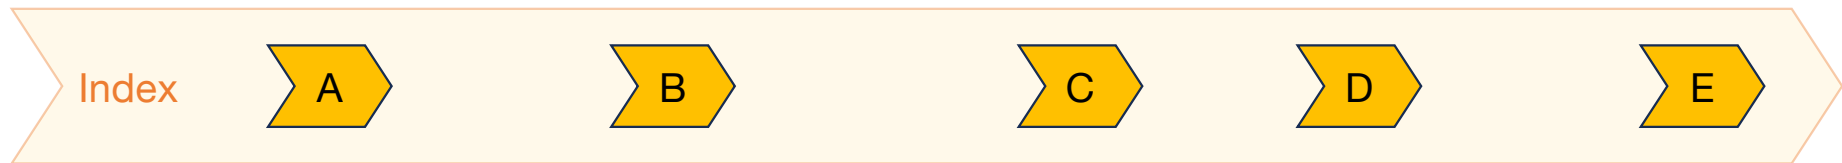
UPDATE a

COMMIT


DEMO



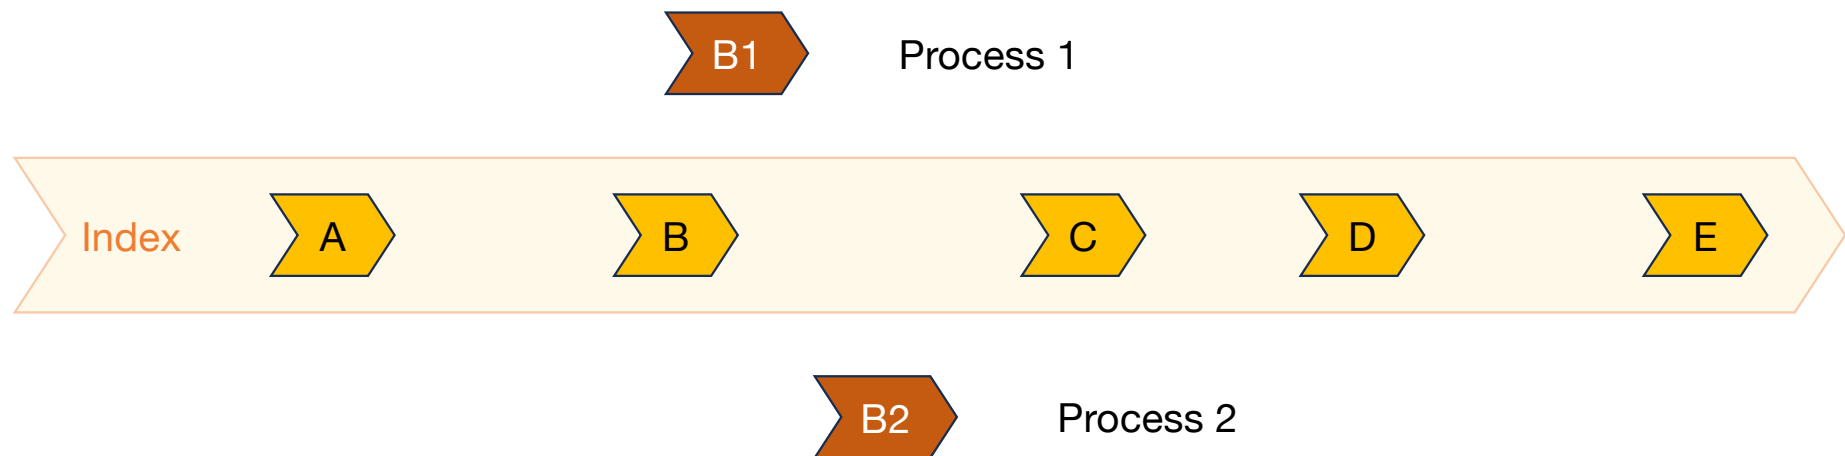
The serializable merge deadlock



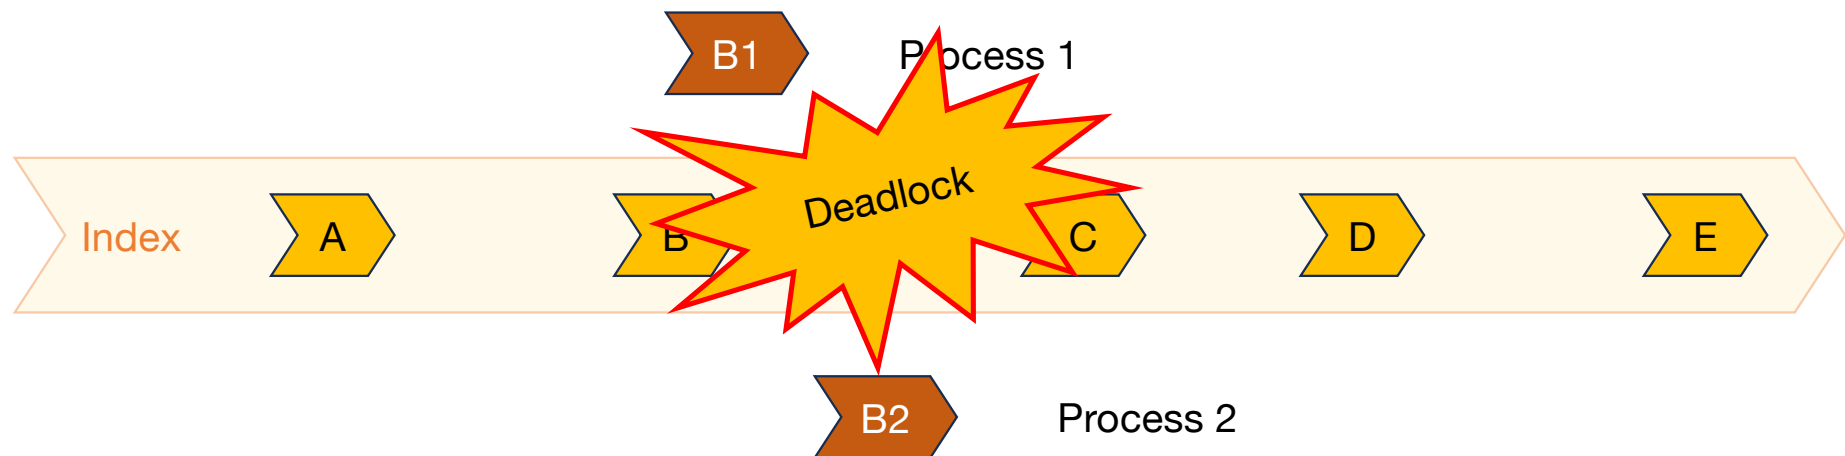
The serializable merge deadlock



The serializable merge deadlock




The serializable merge deadlock



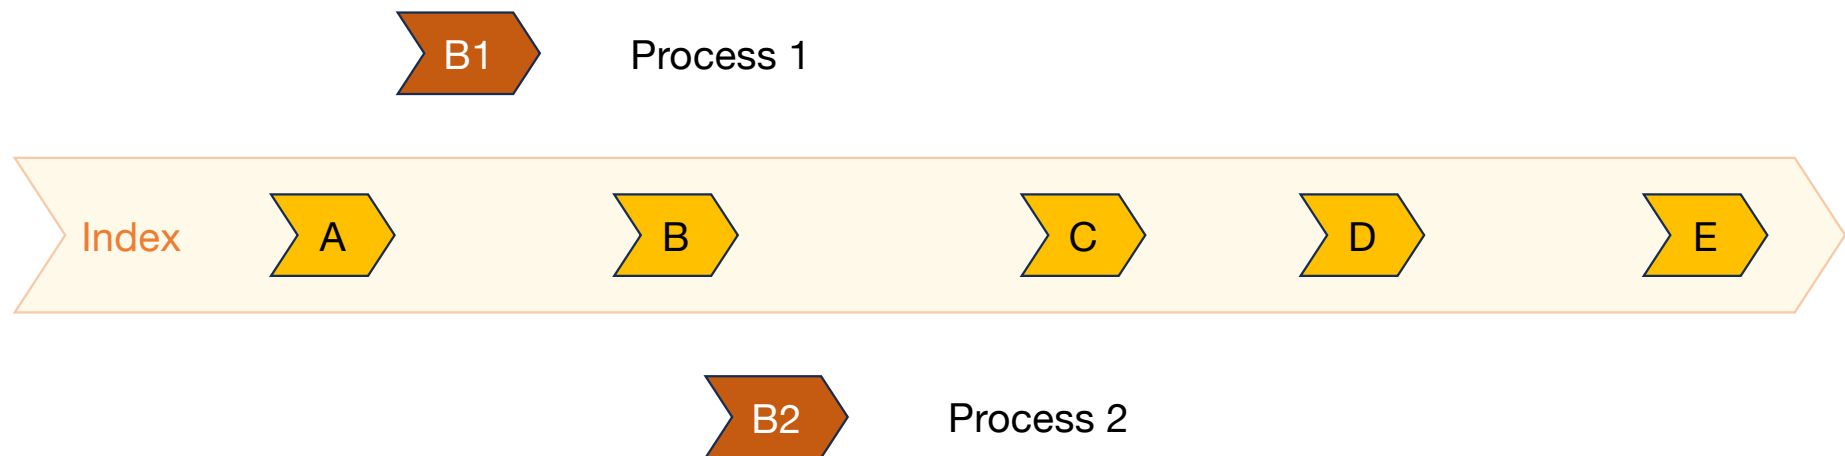
DEMO



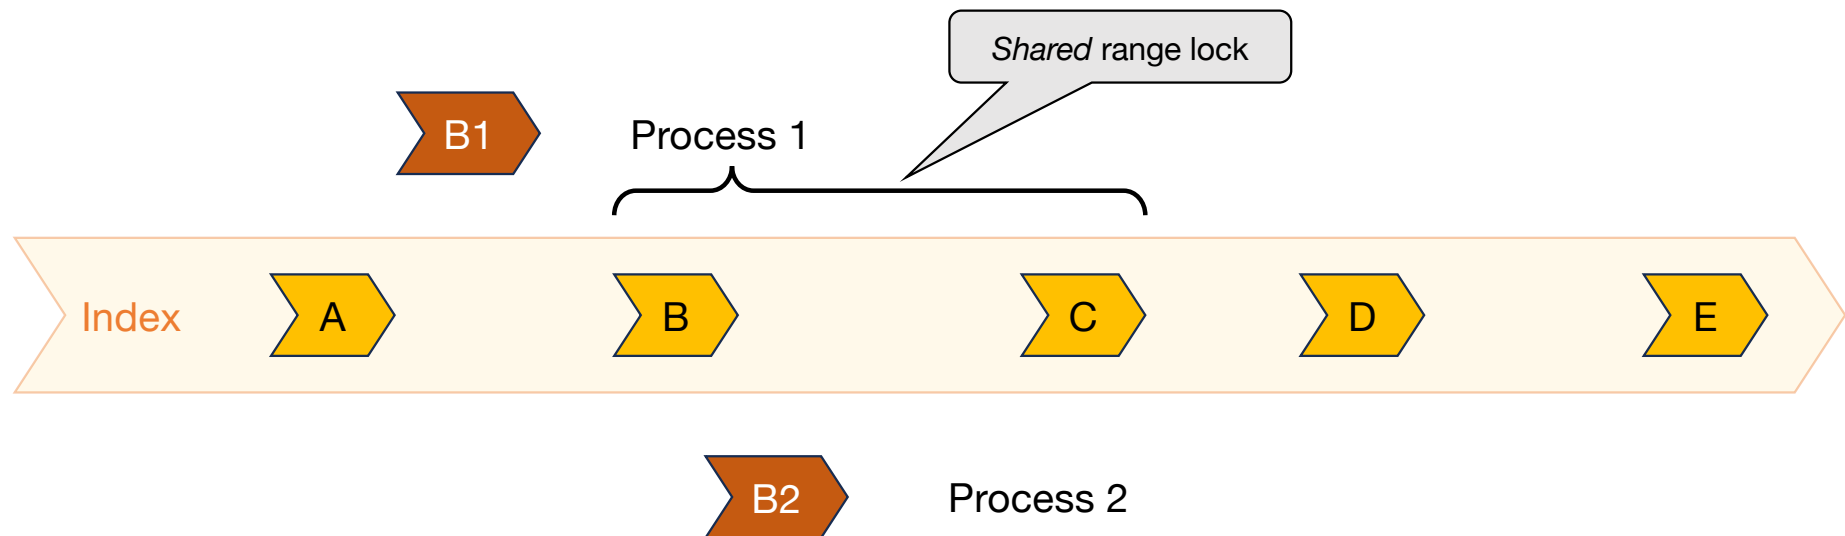
The serializable merge deadlock



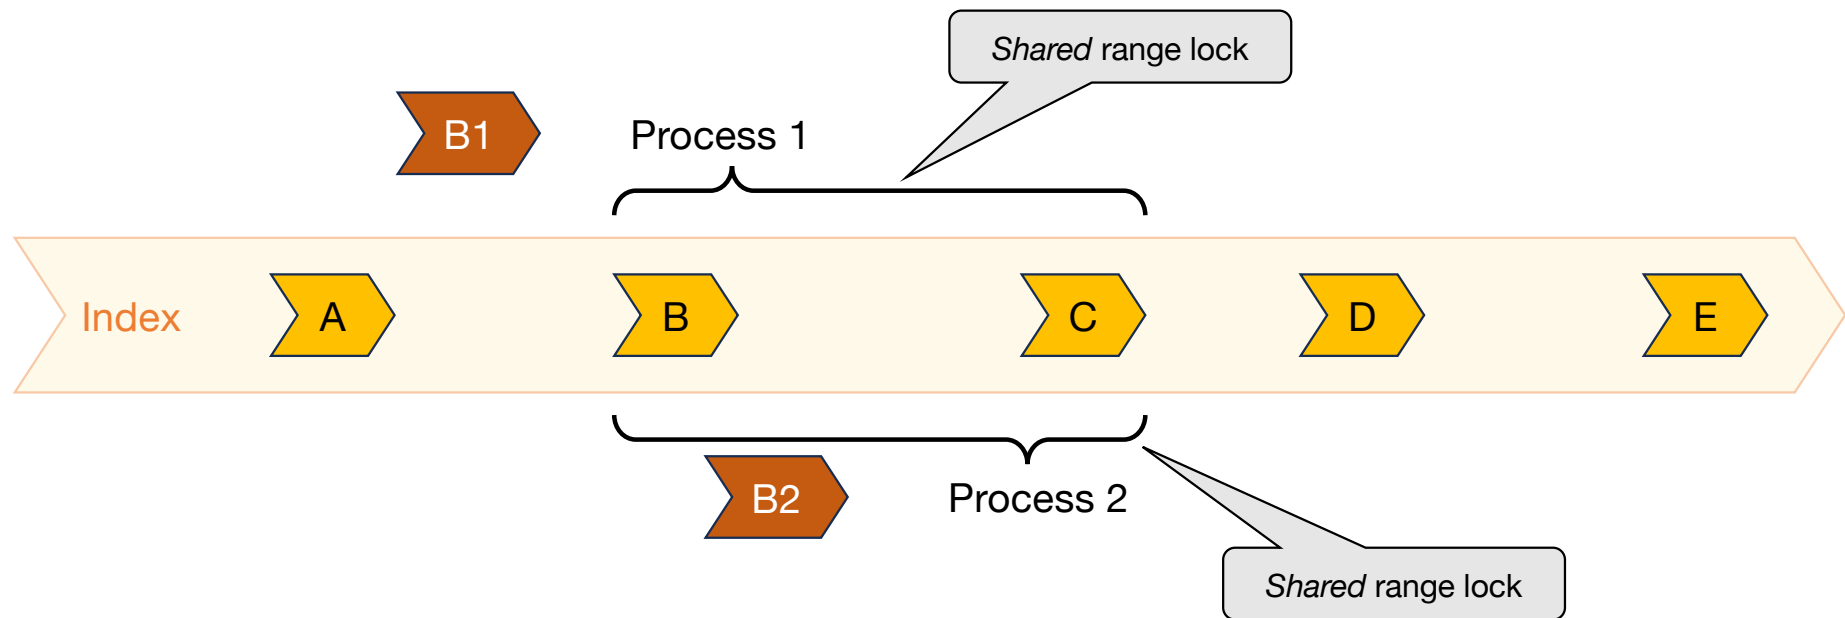
The serializable merge deadlock



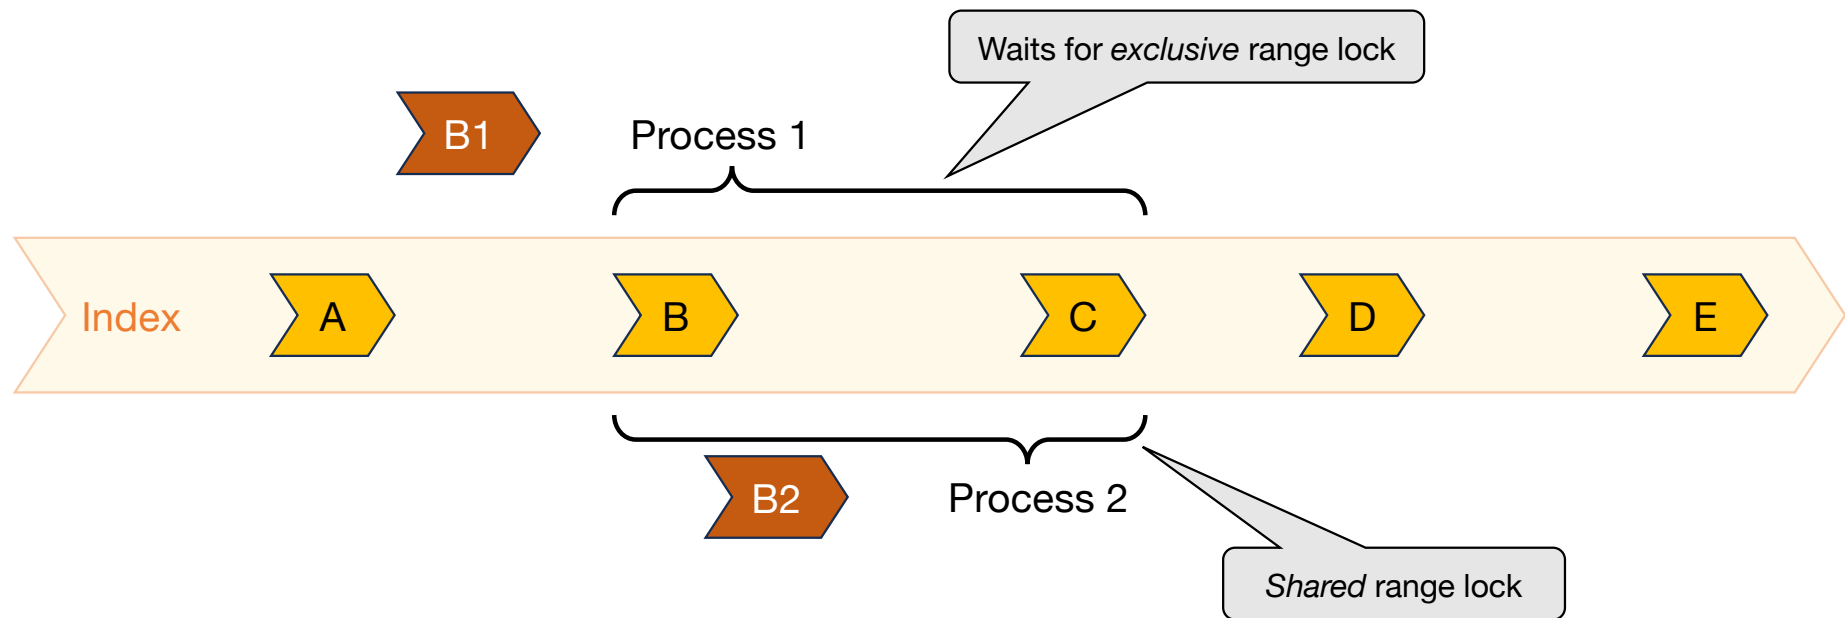
The serializable merge deadlock



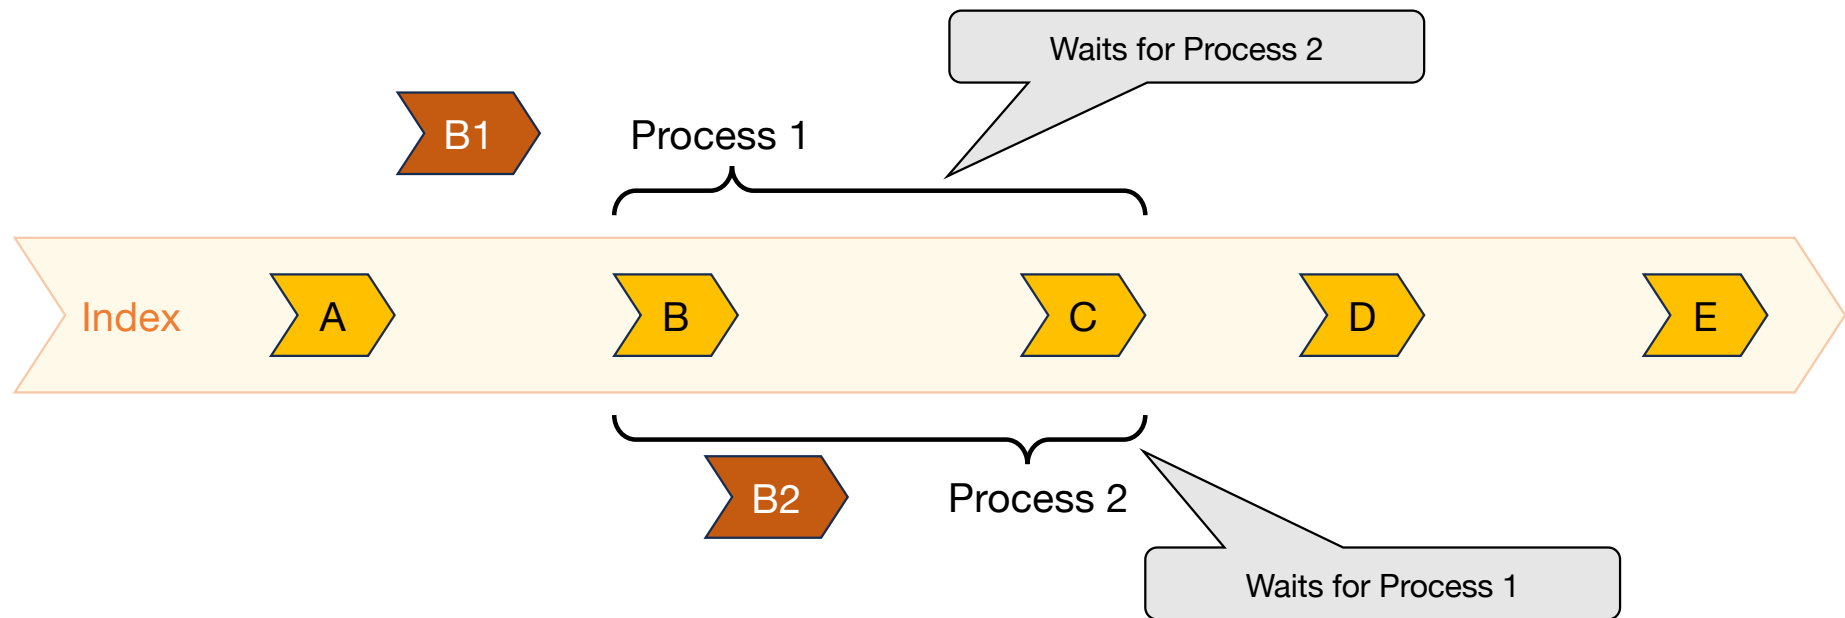
The serializable merge deadlock



The serializable merge deadlock




The serializable merge deadlock




Snapshot isolation



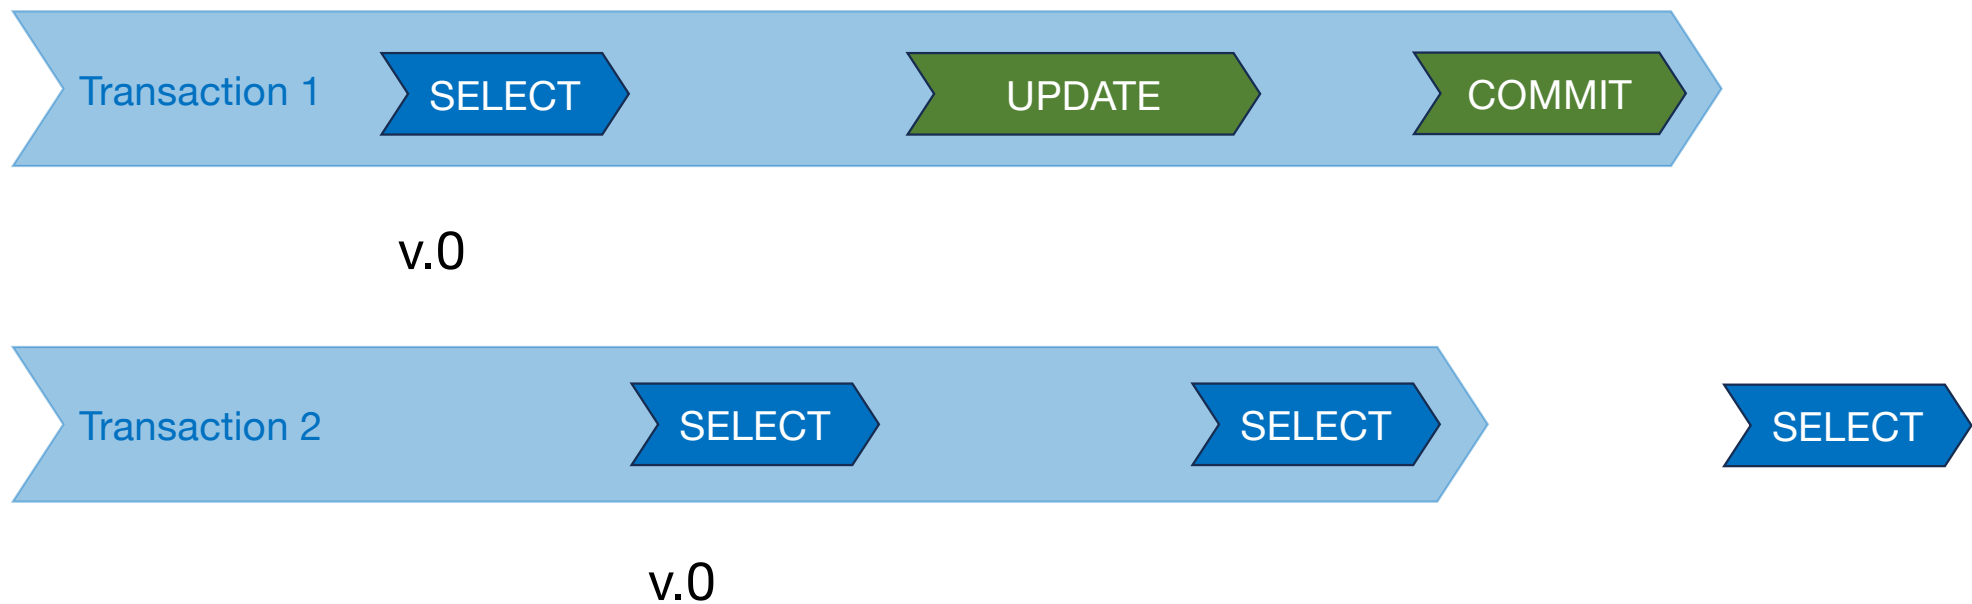
Snapshot isolation



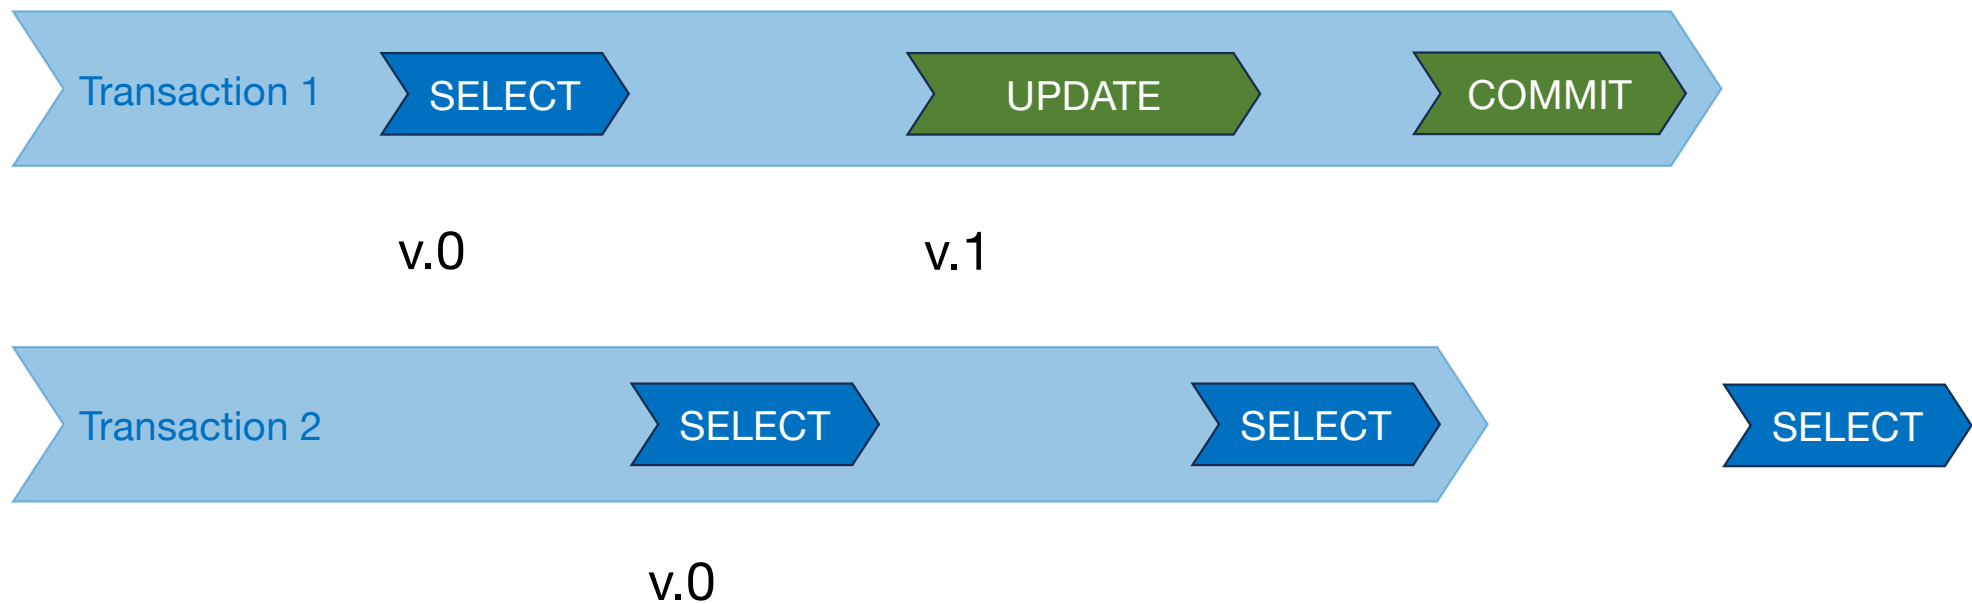
Snapshot isolation



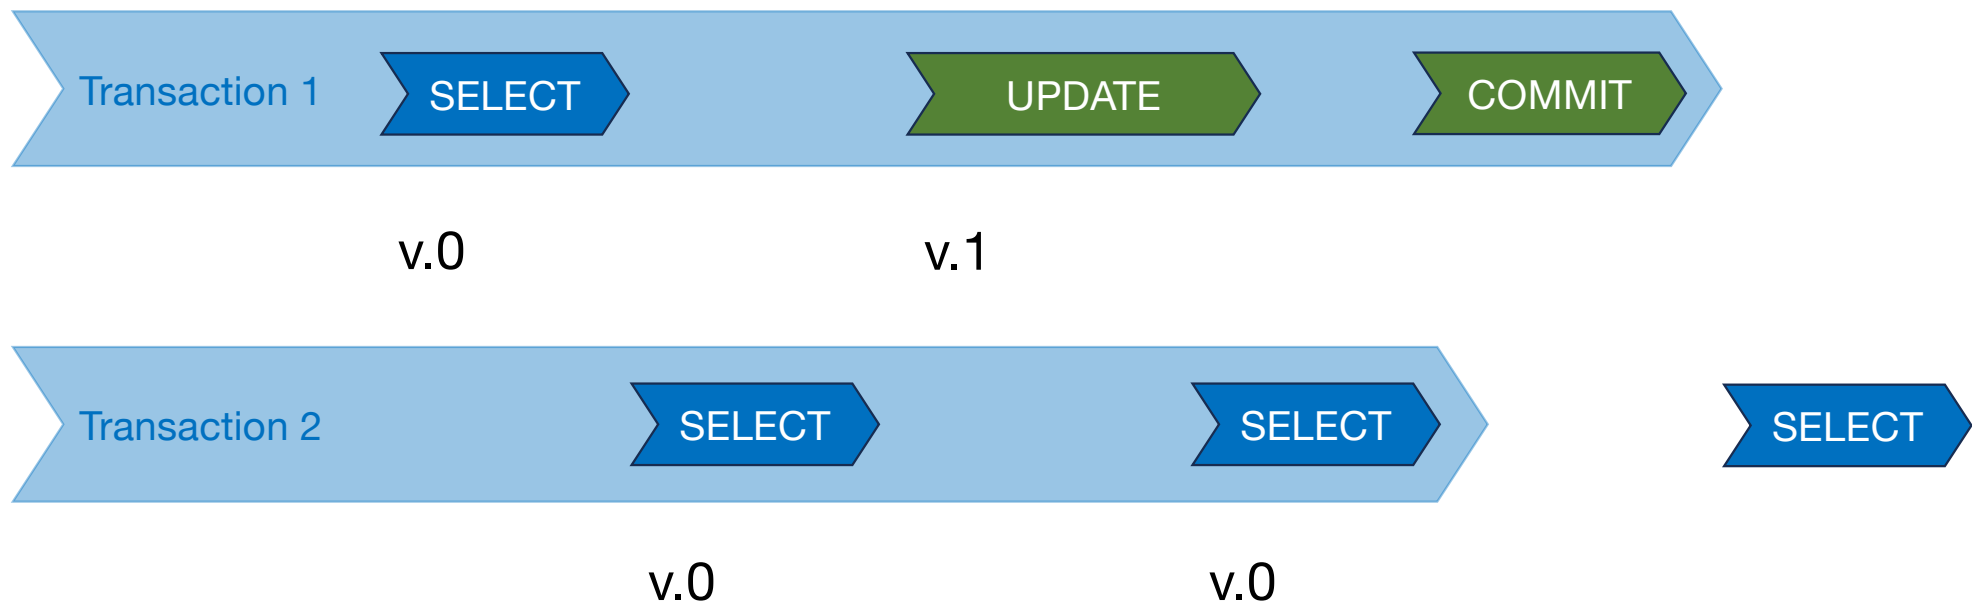
Snapshot isolation



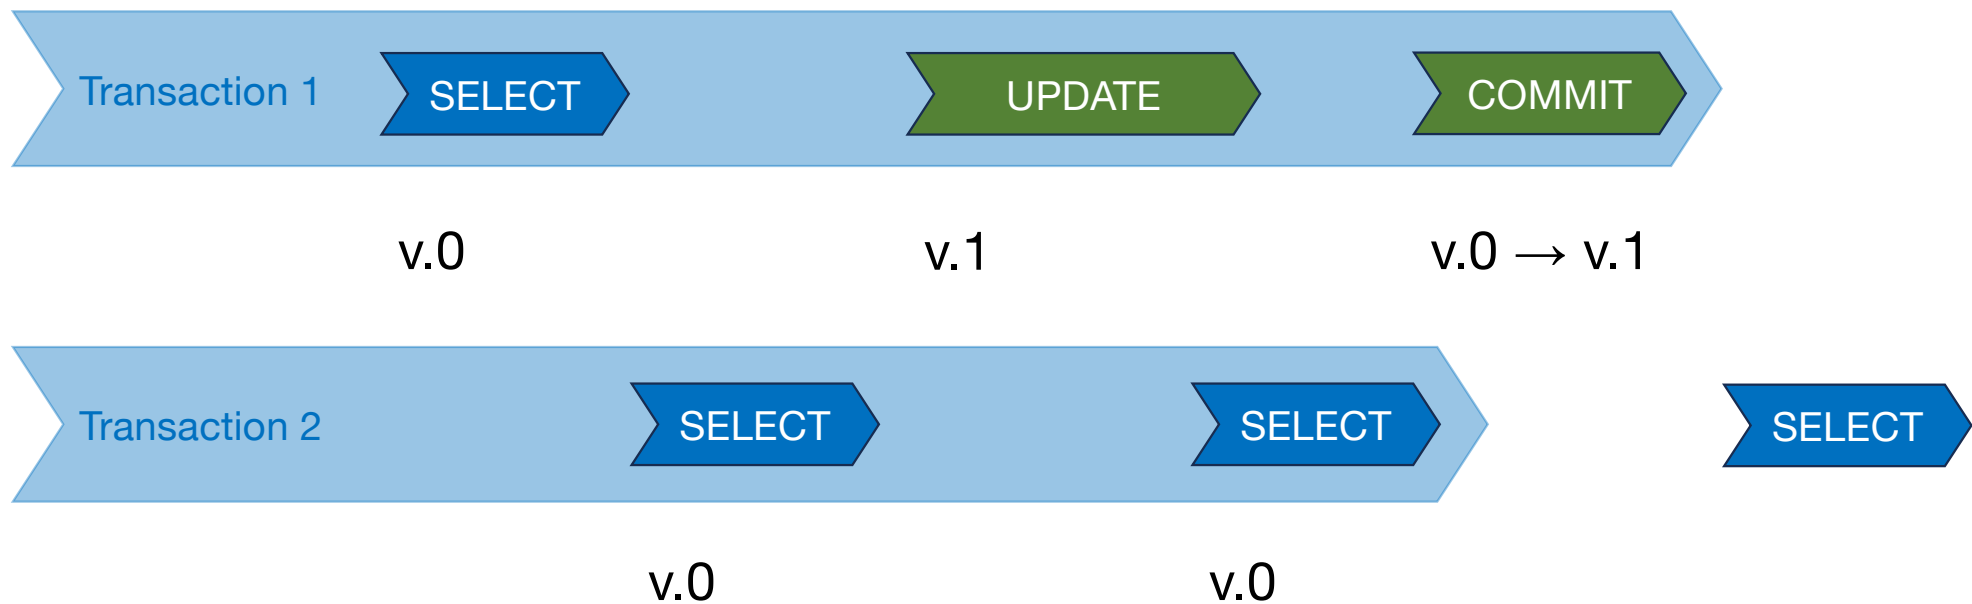
Snapshot isolation




Snapshot isolation



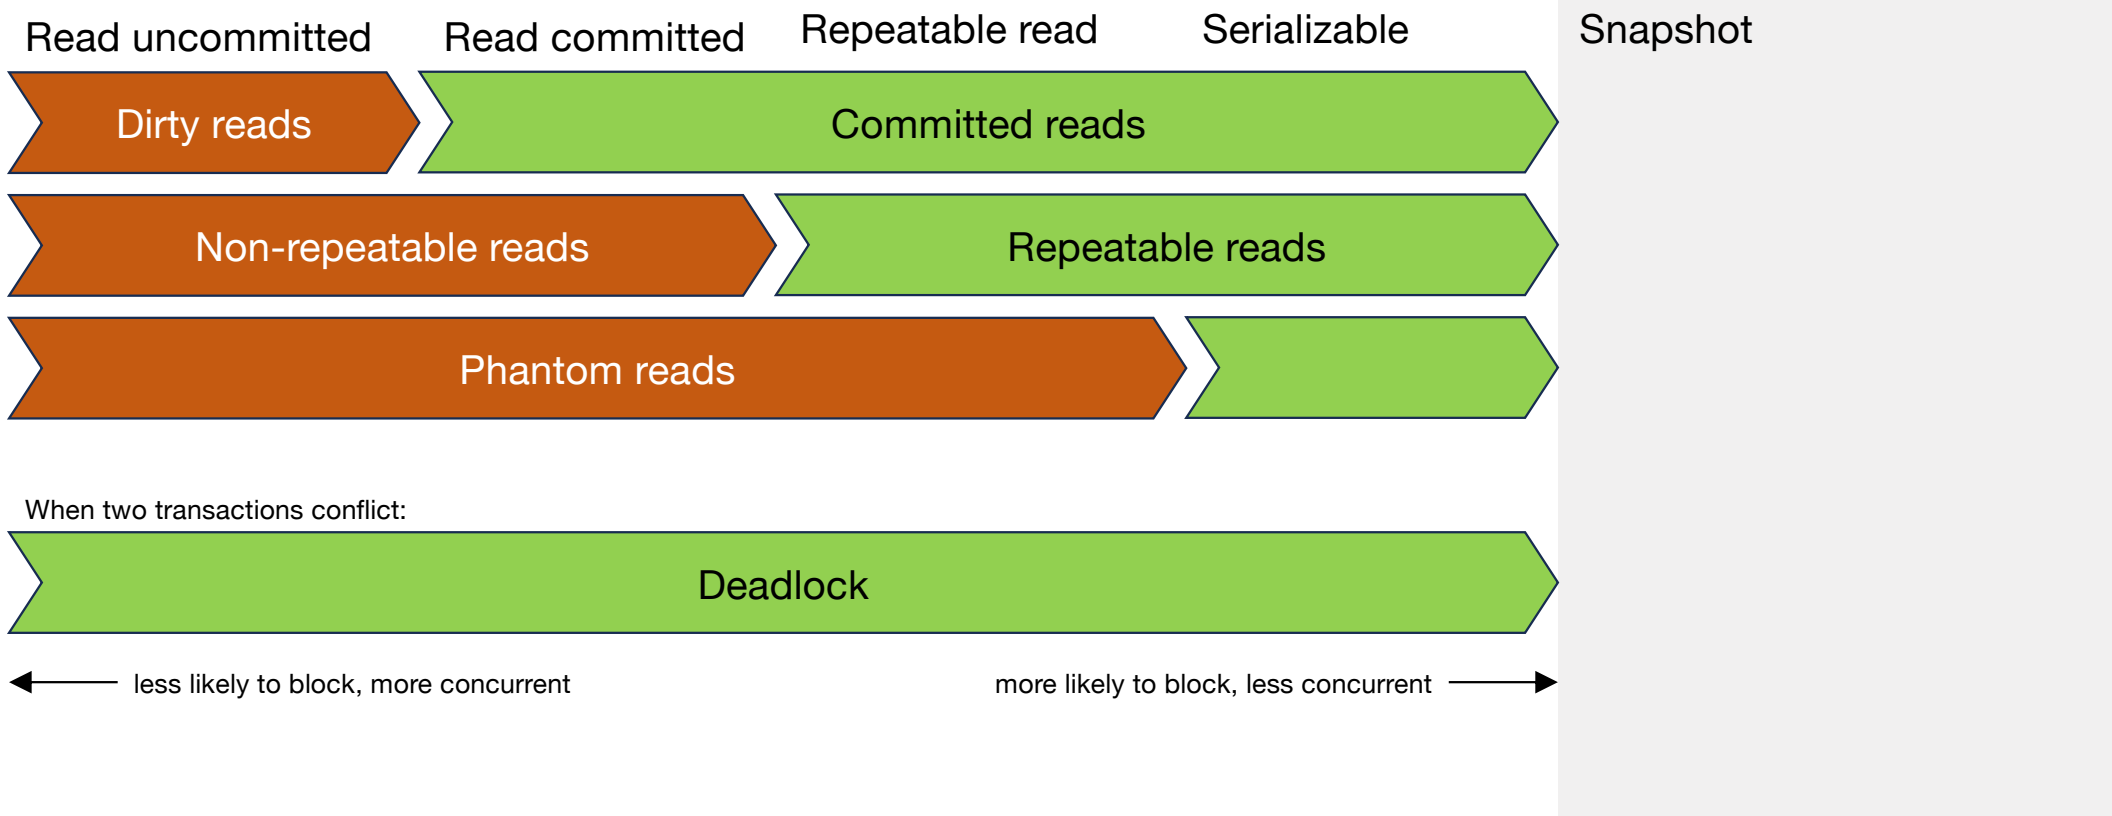
Snapshot isolation



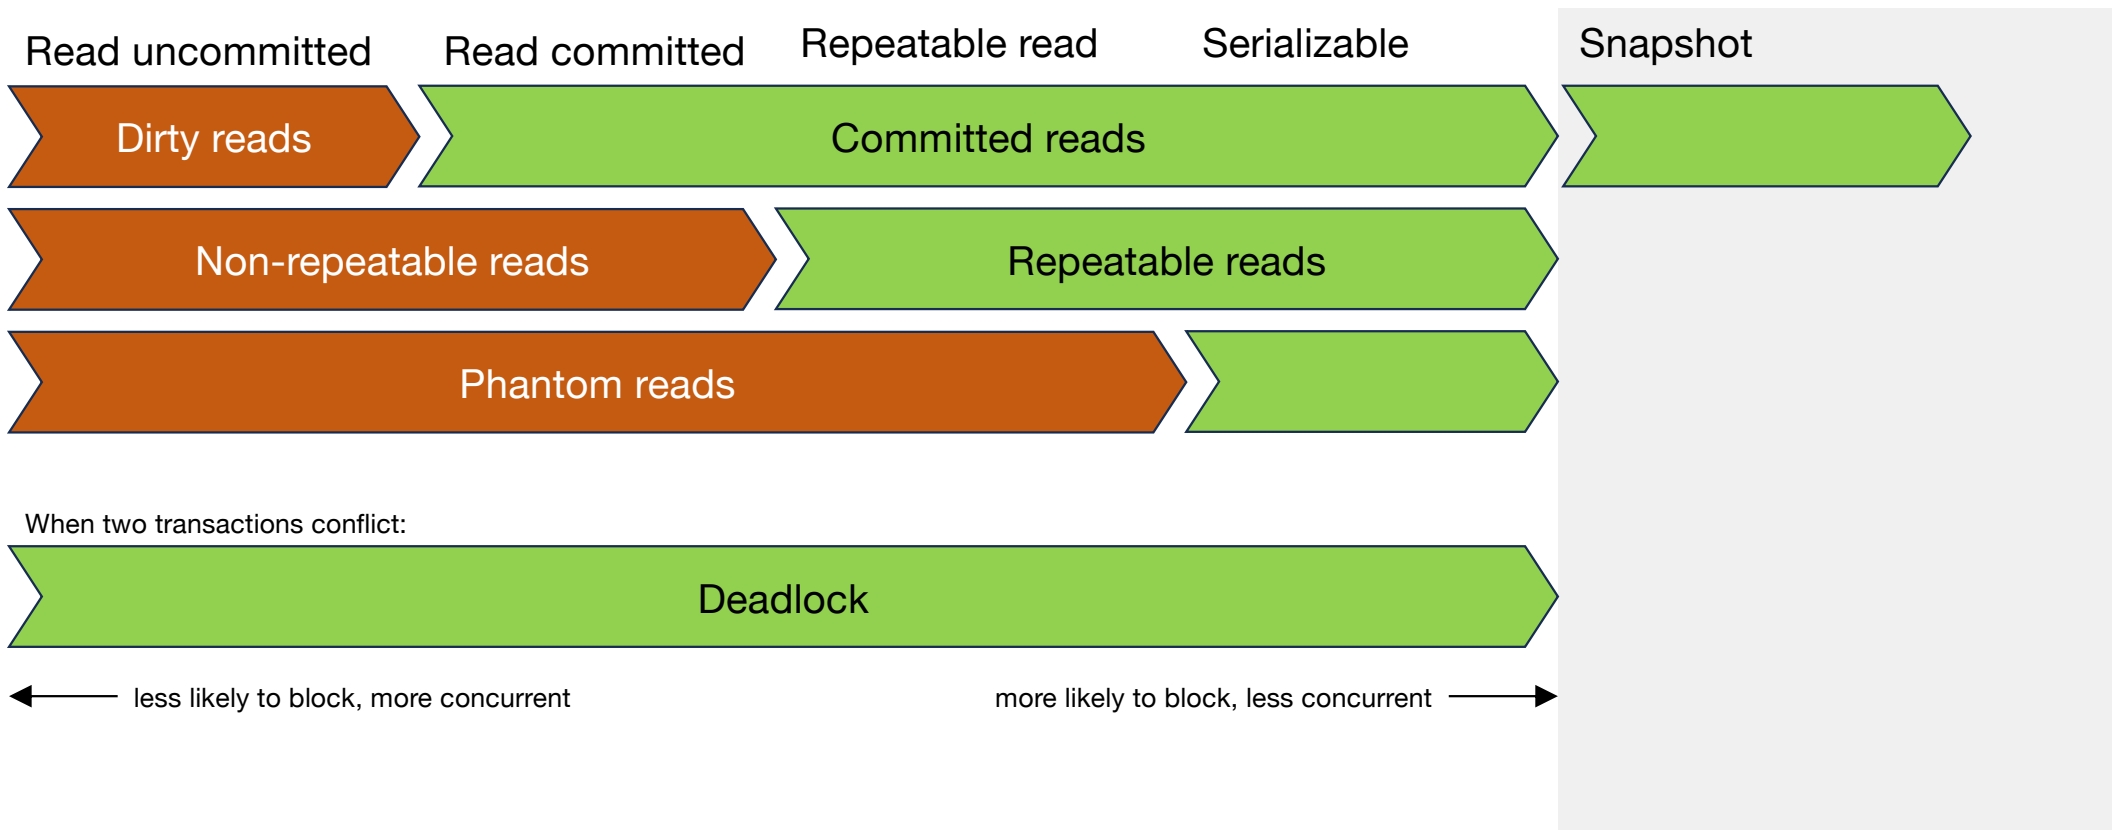
Snapshot isolation



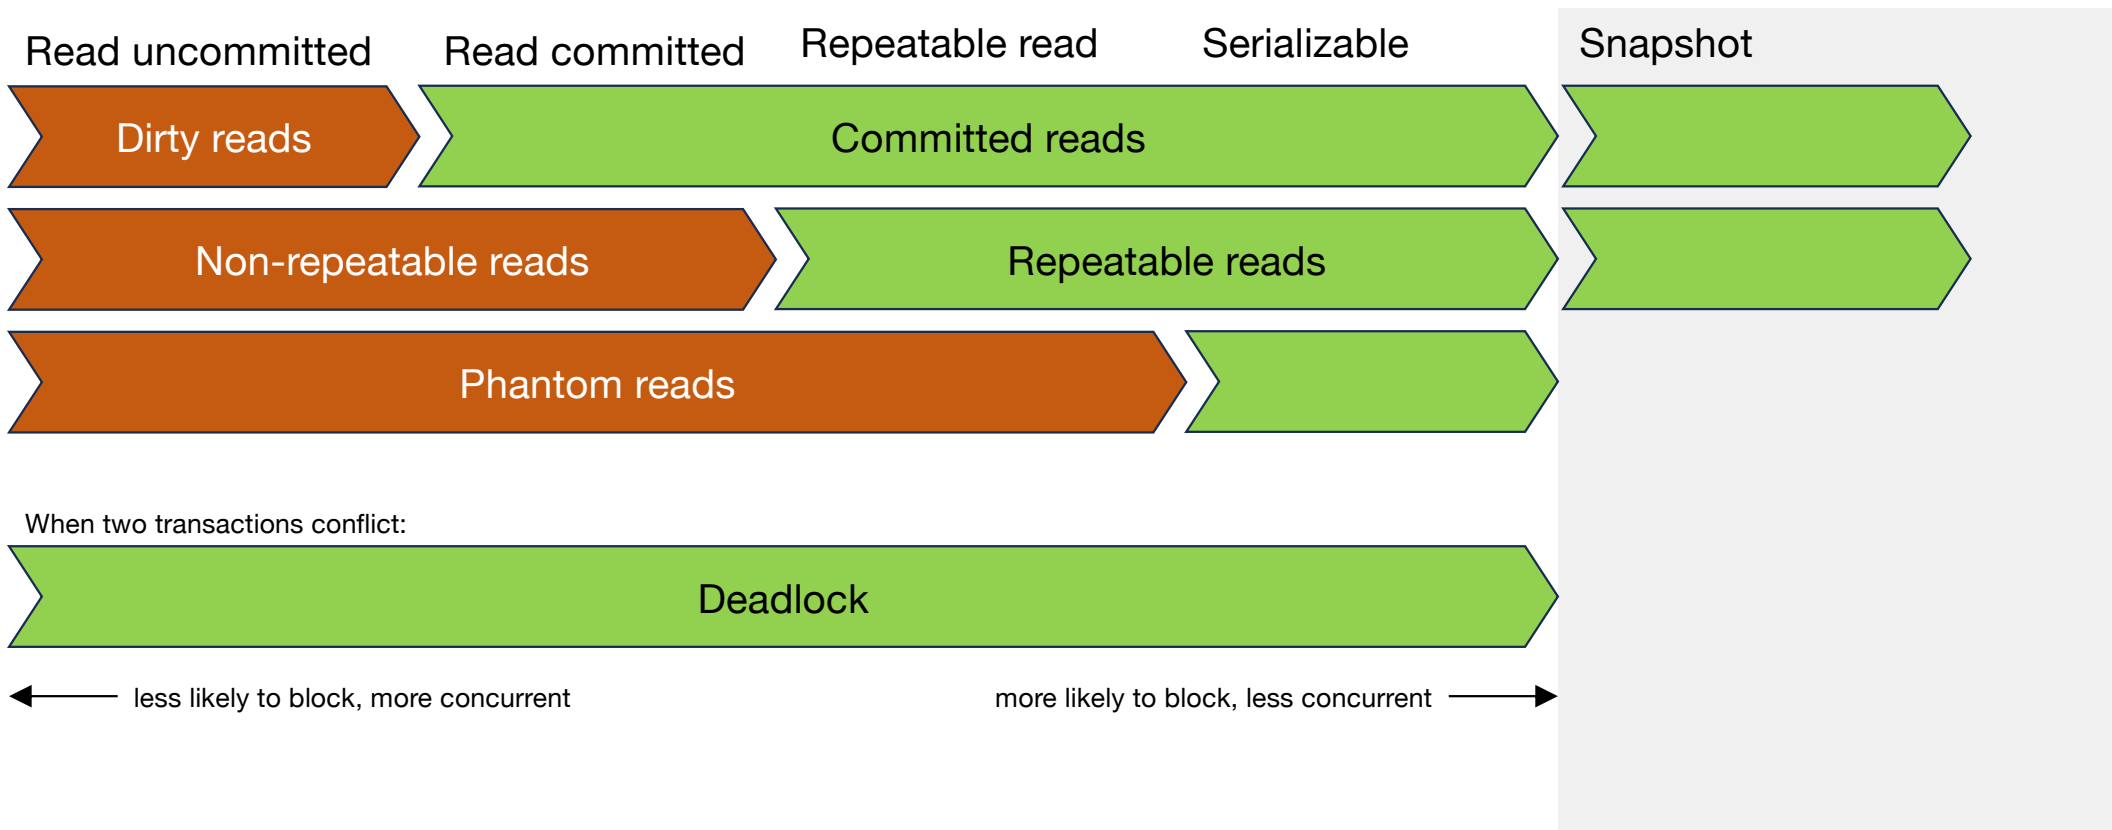
Snapshot isolation



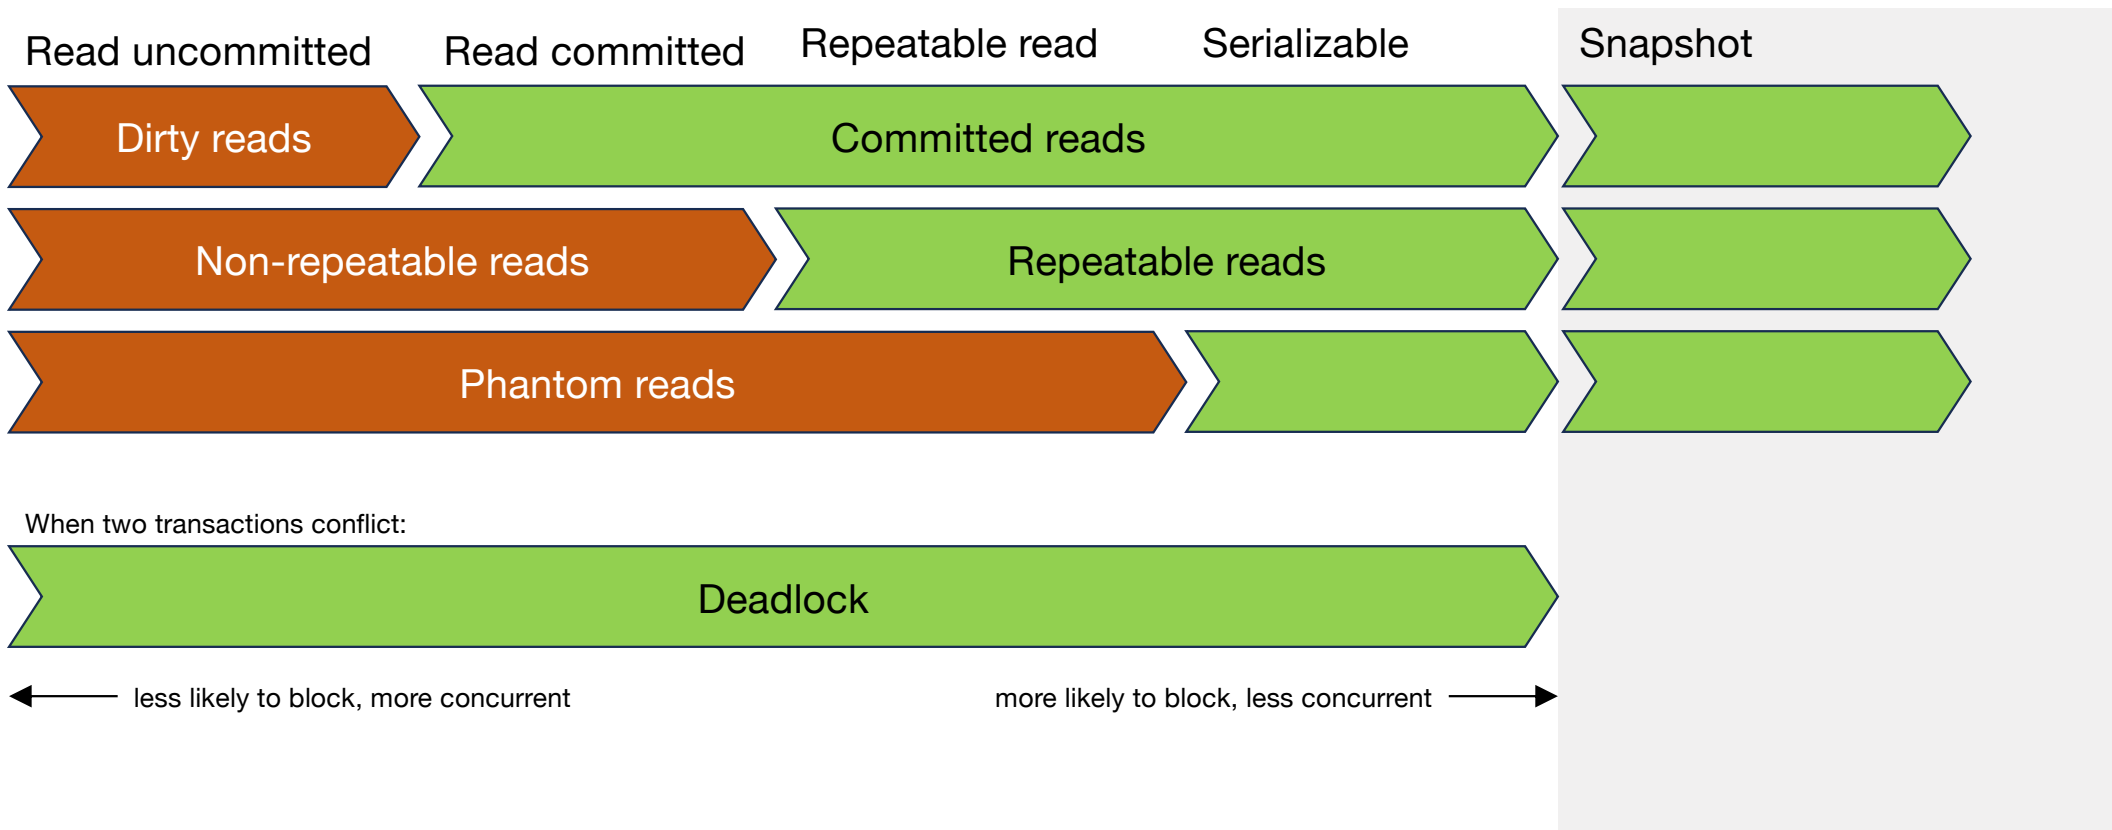
Snapshot isolation



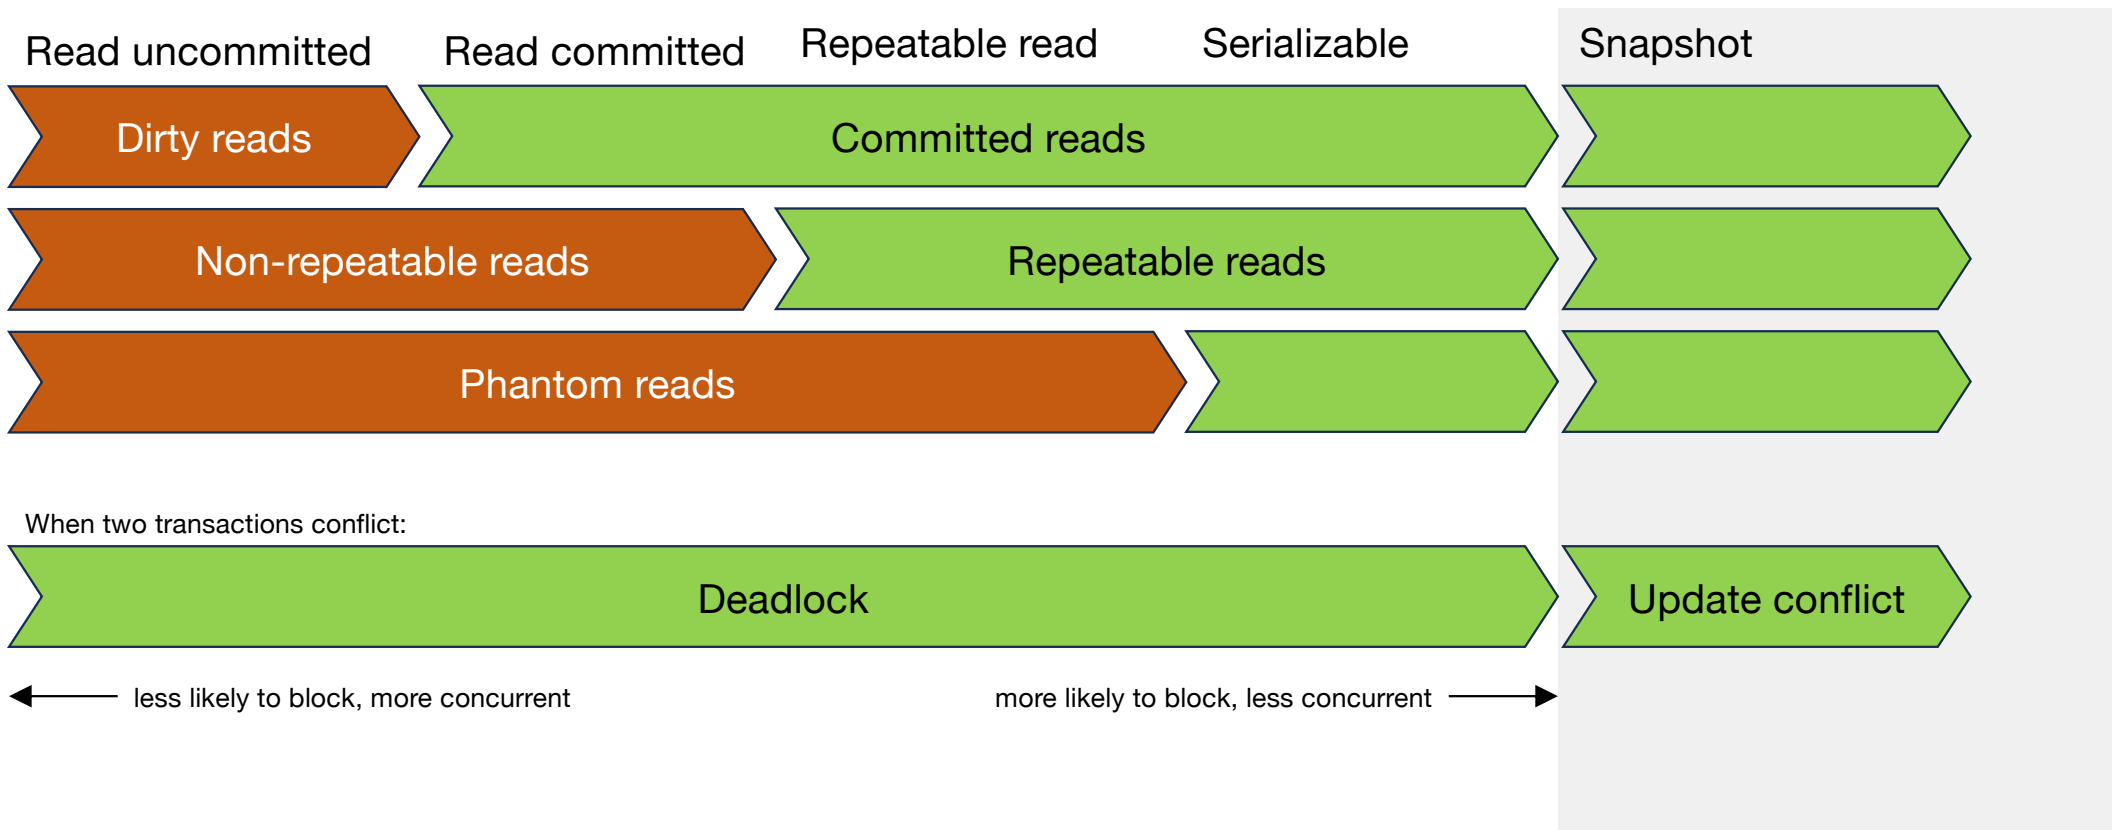
Snapshot isolation



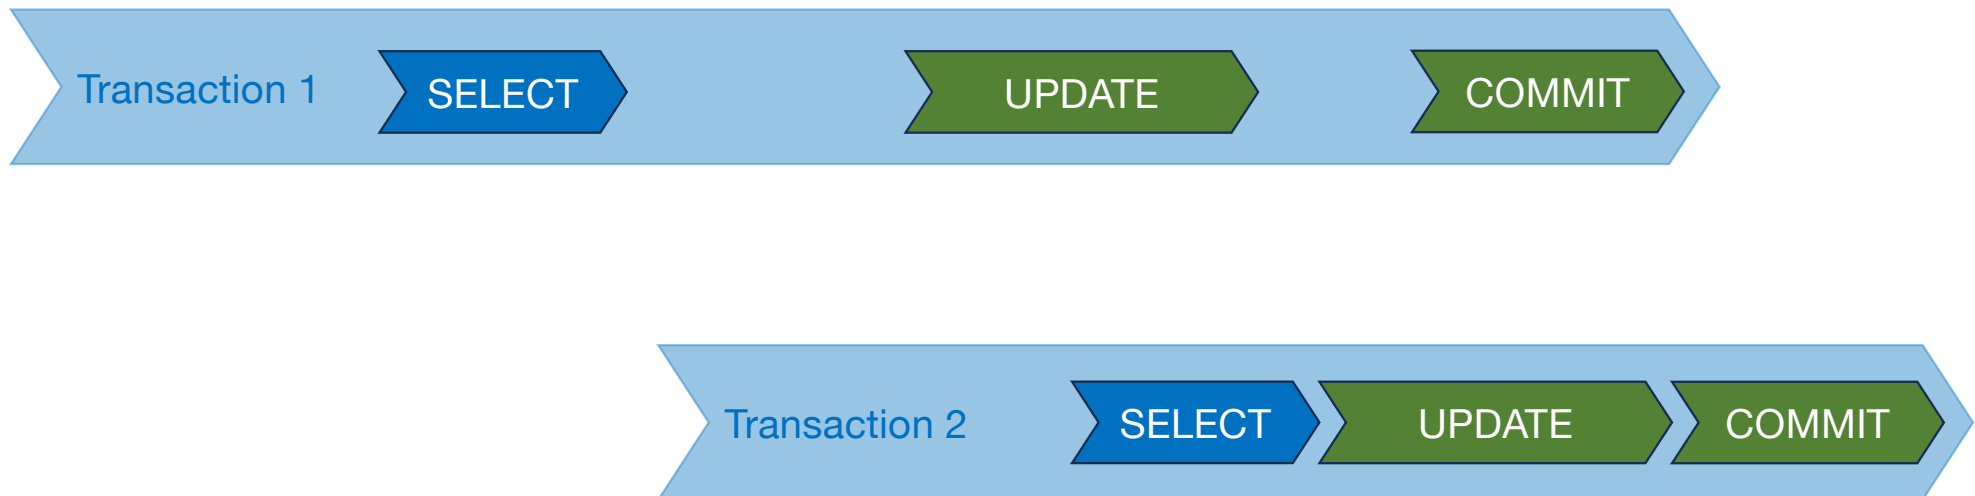
Snapshot isolation



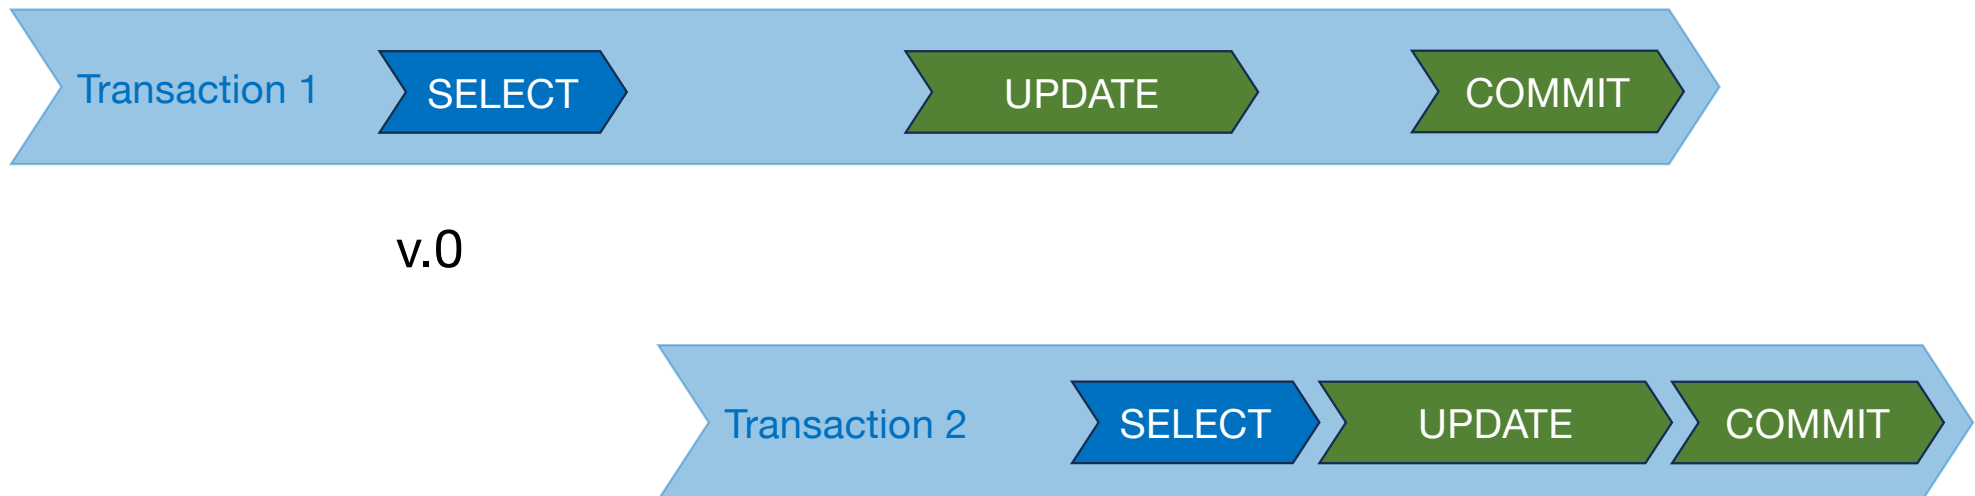
Snapshot isolation



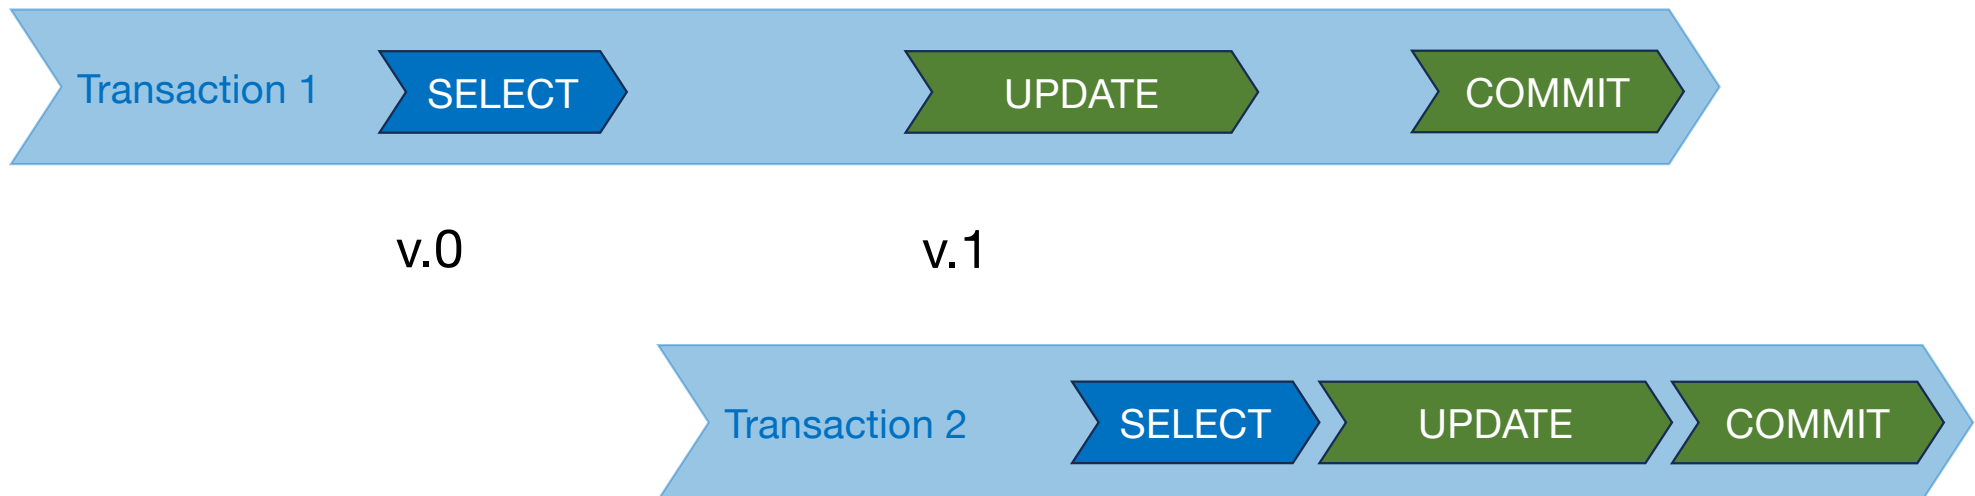
Snapshot isolation: Update conflict



Snapshot isolation: Update conflict



Snapshot isolation: Update conflict



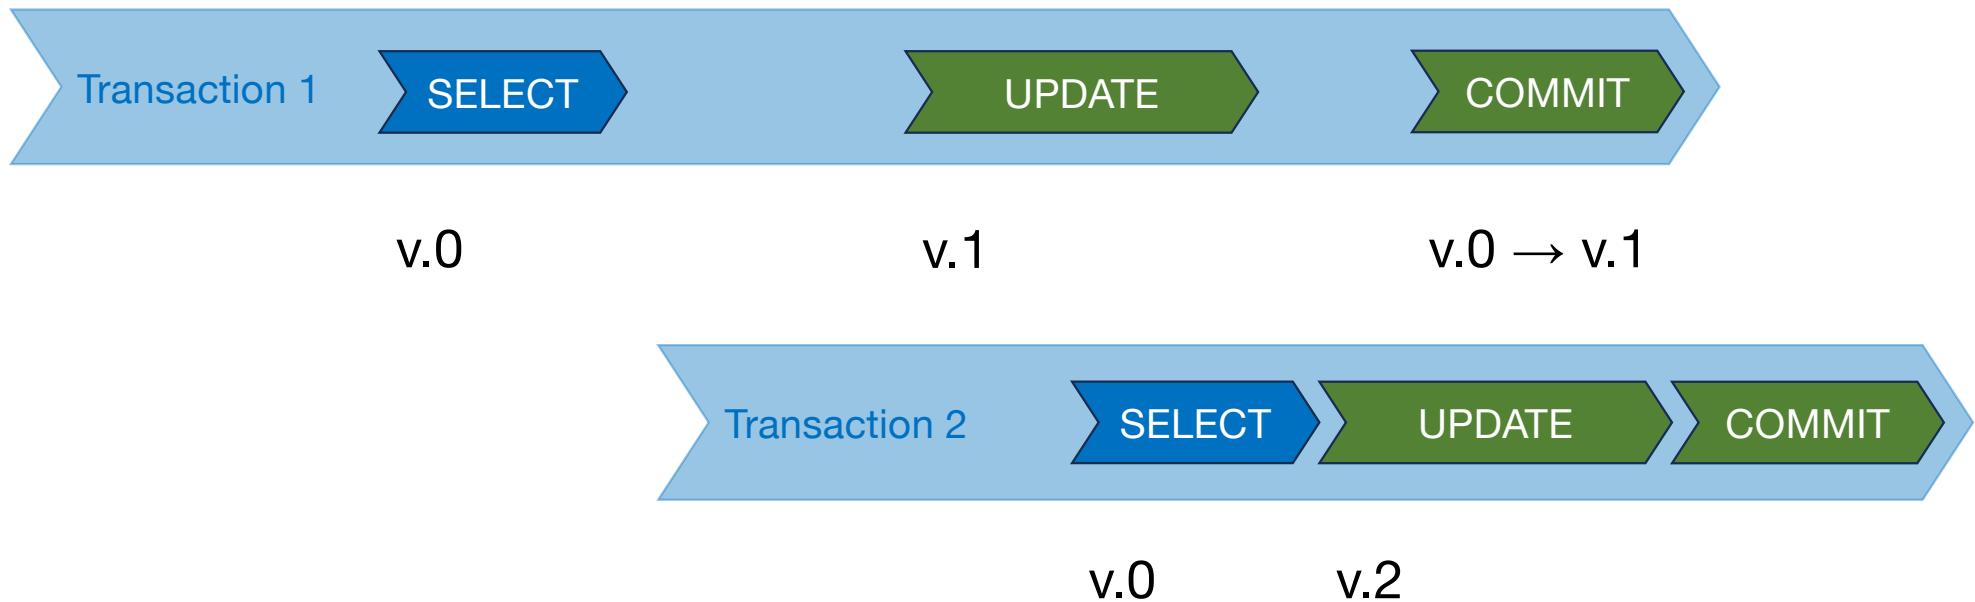
Snapshot isolation: Update conflict



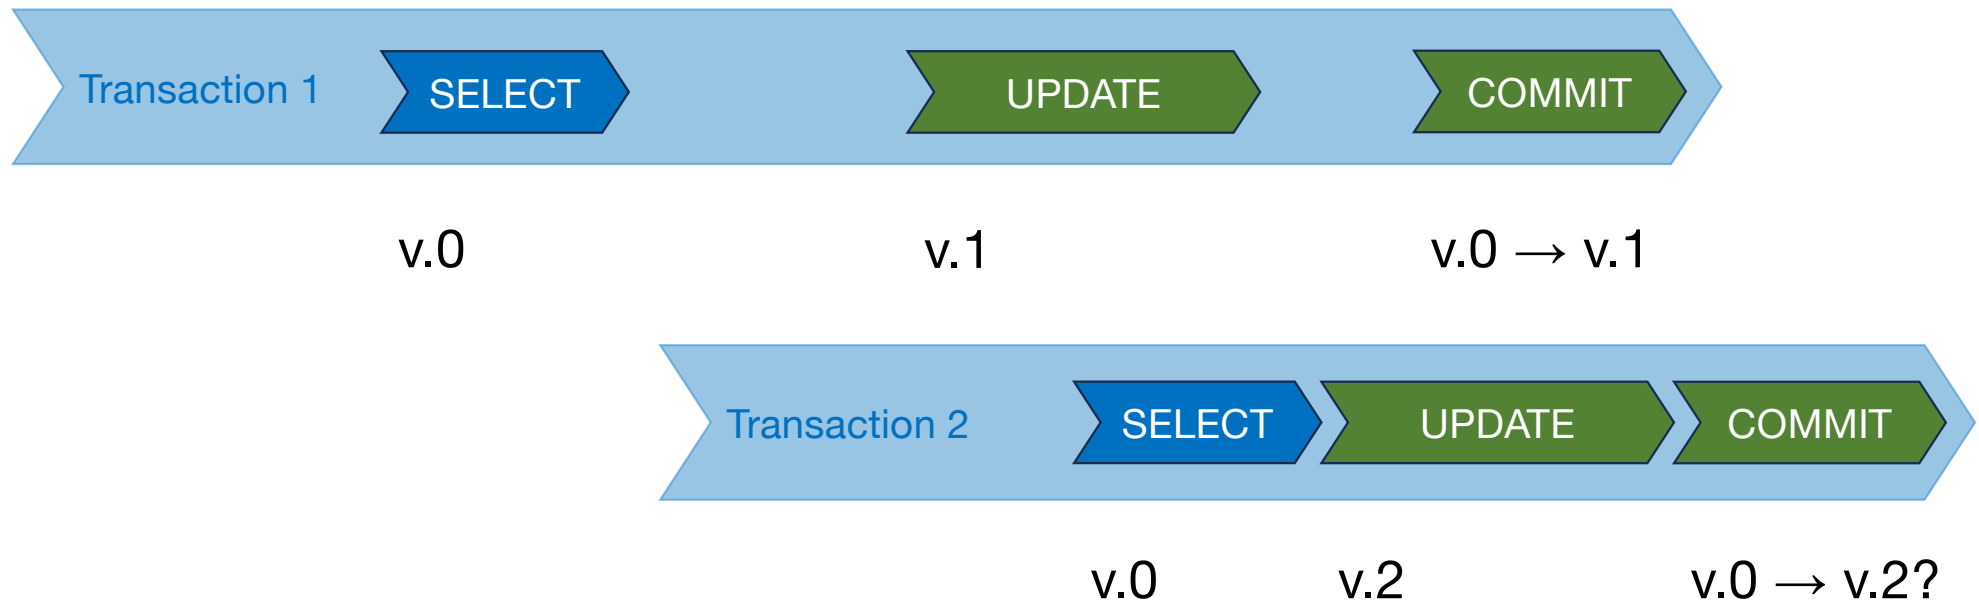
Snapshot isolation: Update conflict



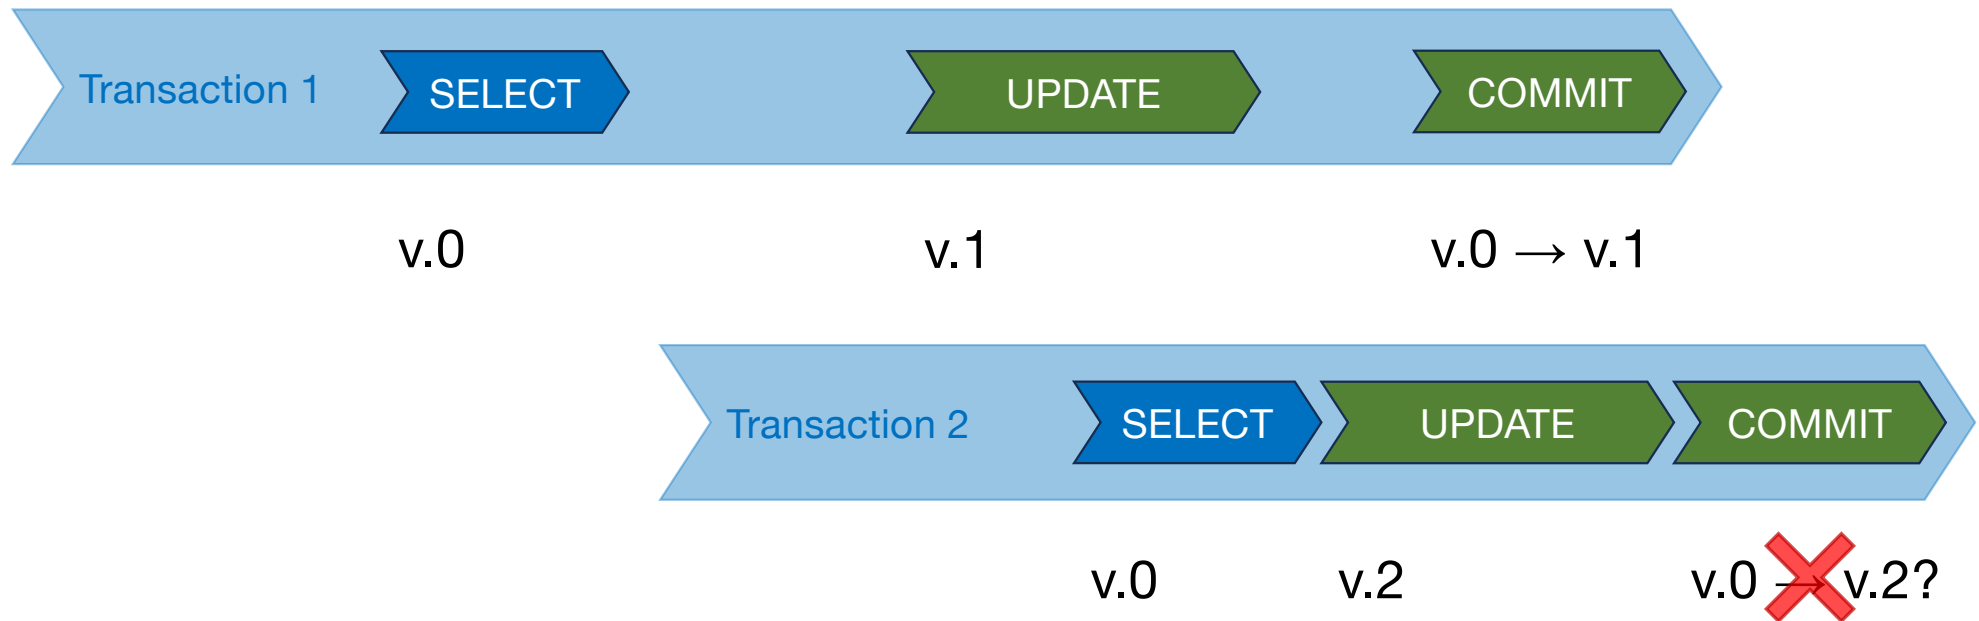
Snapshot isolation: Update conflict



Snapshot isolation: Update conflict



Snapshot isolation: Update conflict



DEMO



RCSI vs. Snapshot

- Snapshot Isolation is set per *transaction*.
- Snapshot Isolation protects the *transaction*.
- Snapshot Isolation requires code change.
- Conflicting writes will cause *update conflict*.

- Read Committed Snapshot Isolation is a *database* setting.
- Read Committed Snapshot Isolation protects the *statement*.
- Conflicting writes will cause *blocking*.

- Both require testing, because they behave differently.


More details: <https://brentozar.com/go/rcsi>




How the transaction log works




How the transaction log works




How the transaction log works




How the transaction log works



How the transaction log works



How the transaction log works



How the transaction log works



Durability



SQL

KONFERENZ

2024

Write to
data file

Durability



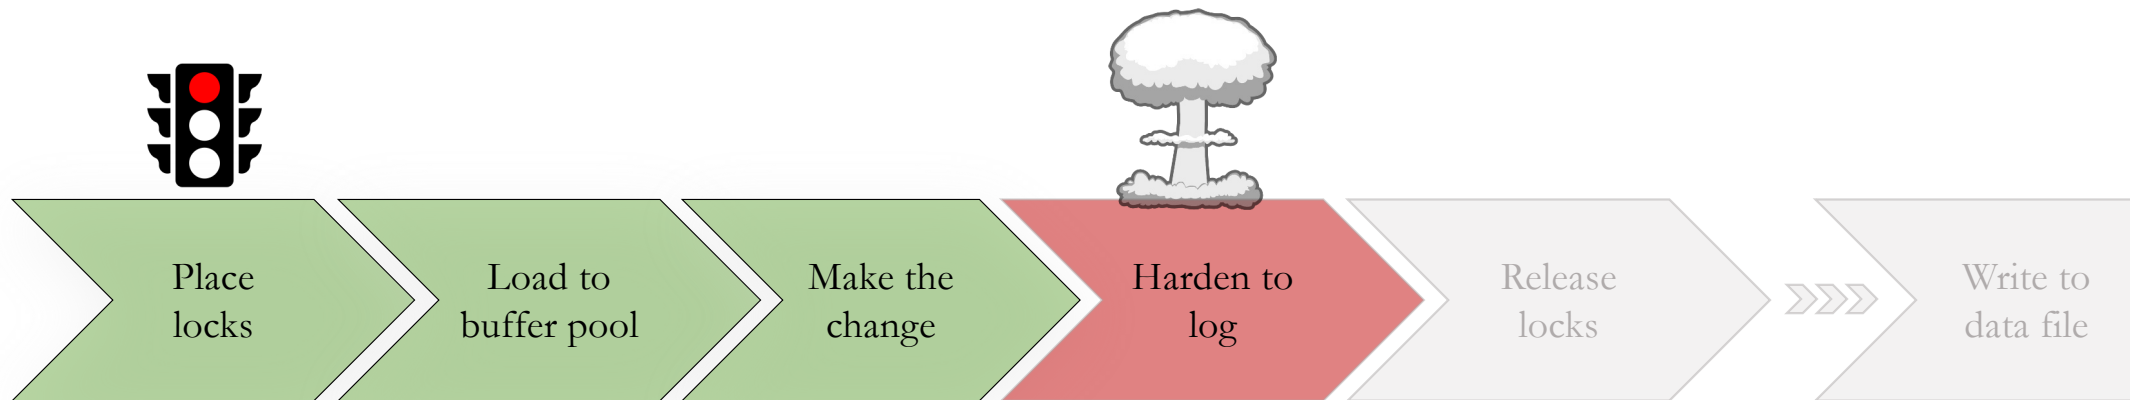
Durability



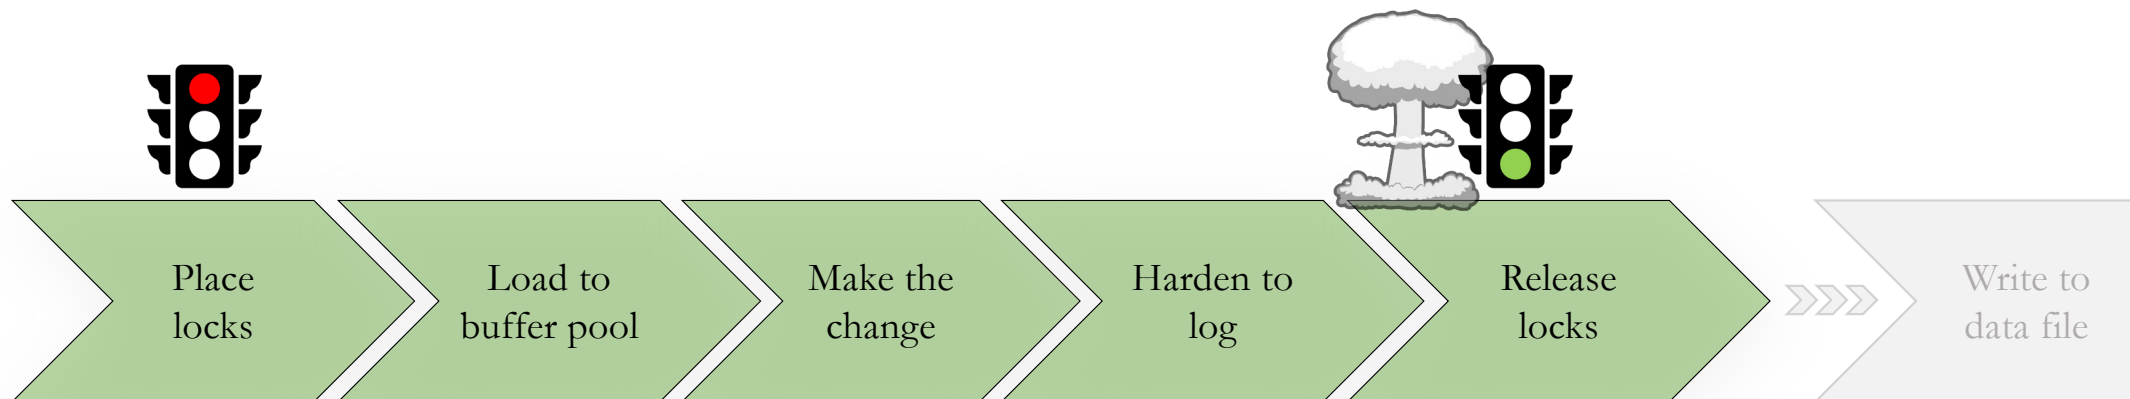
Durability



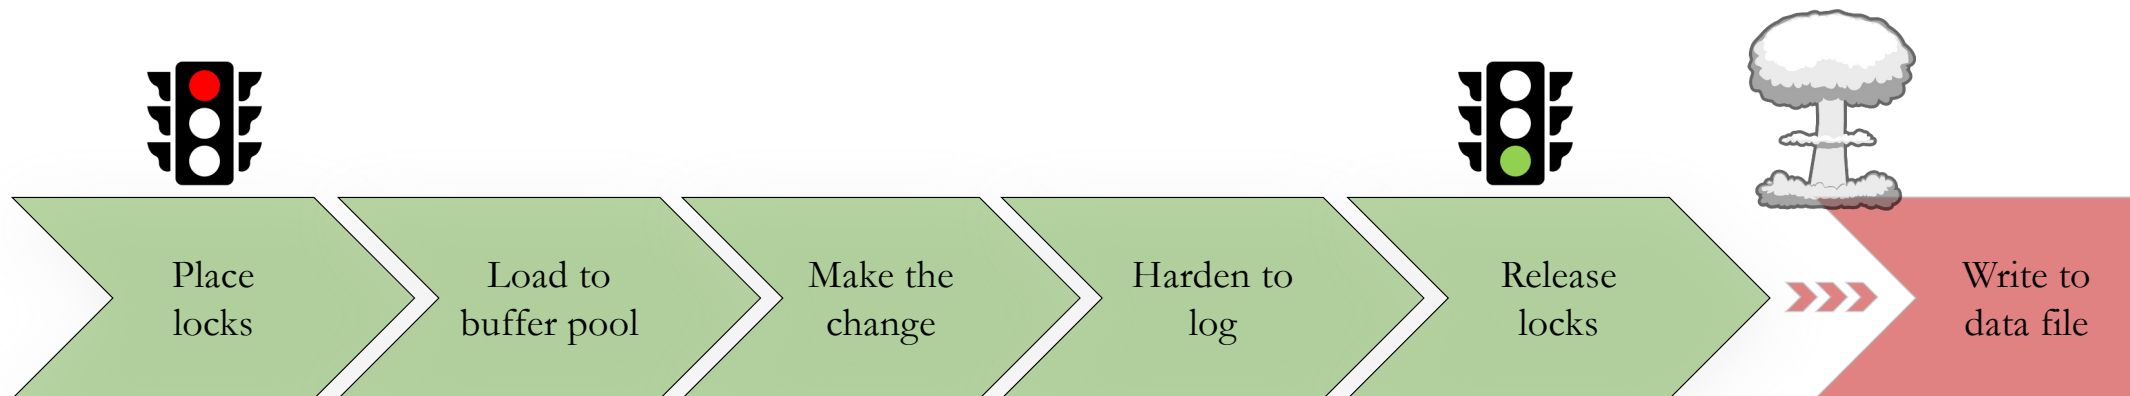
Durability



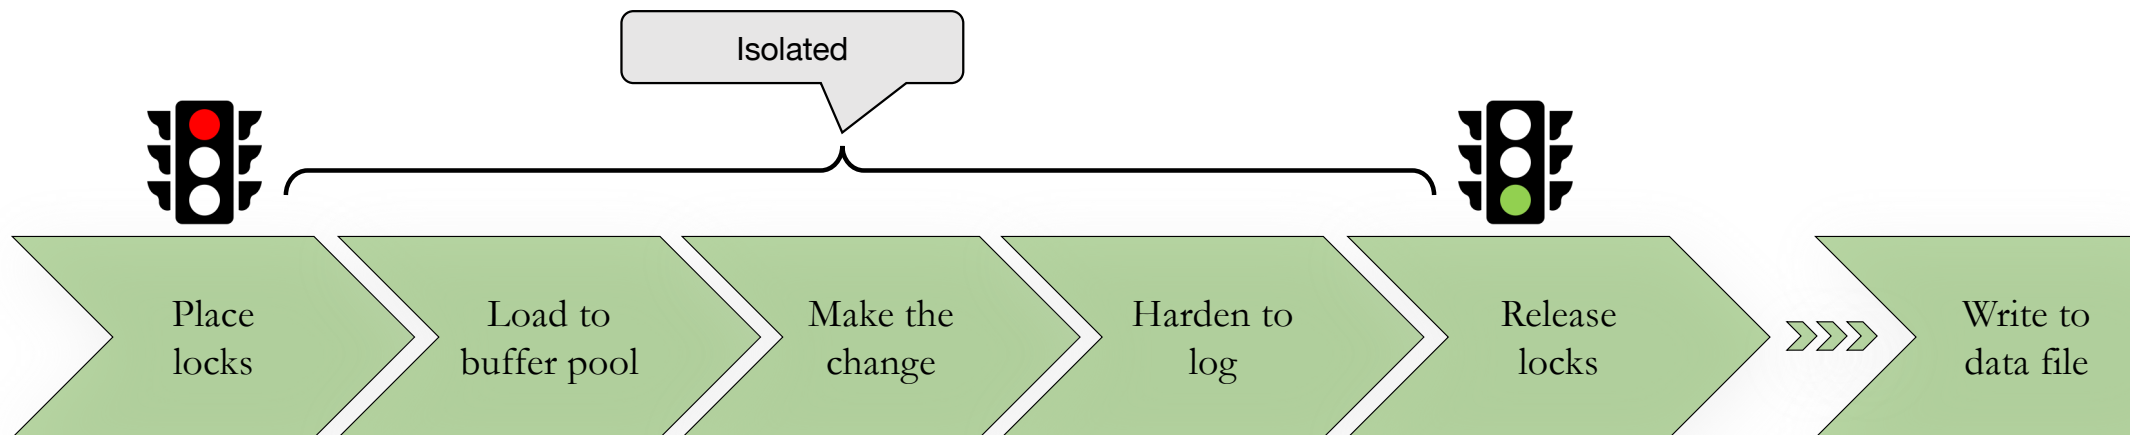
Durability



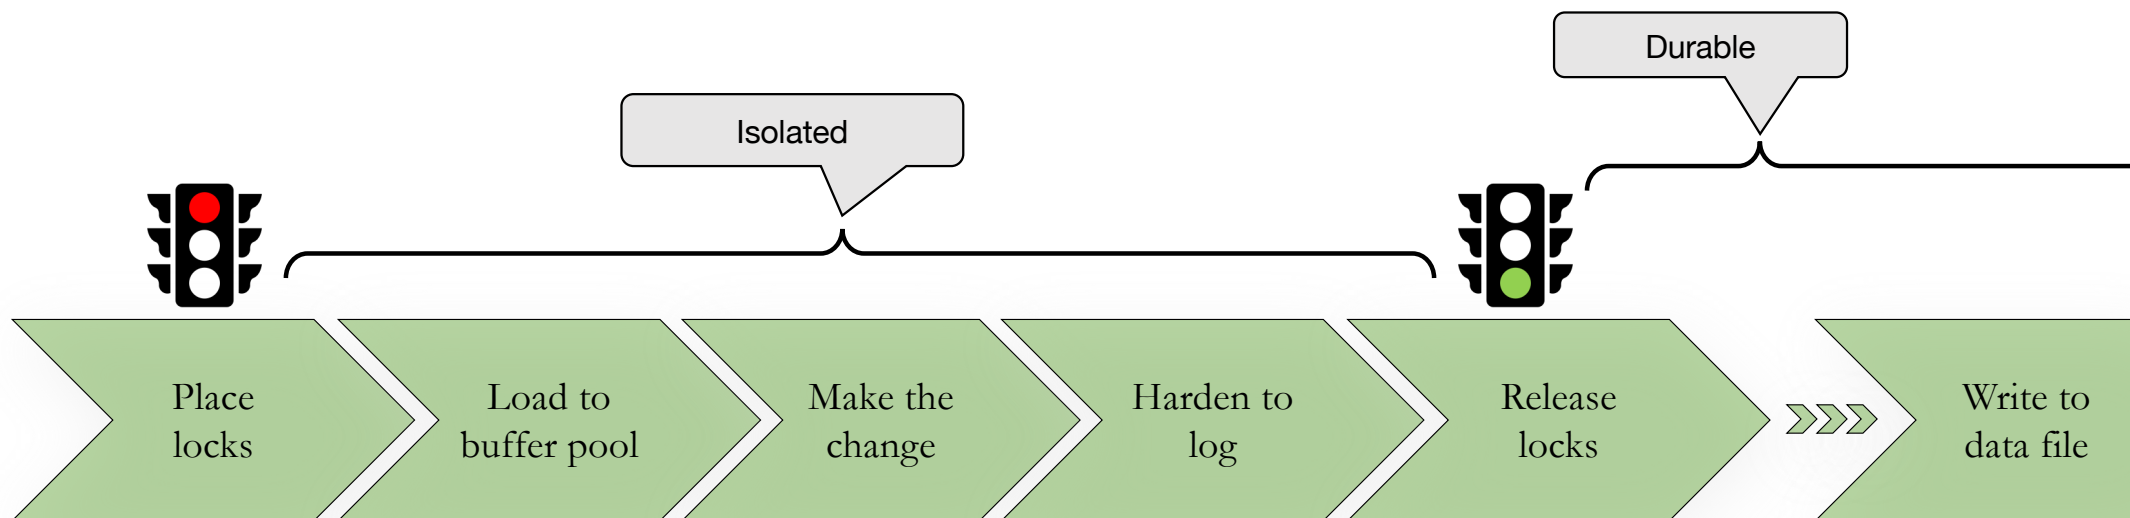
Durability



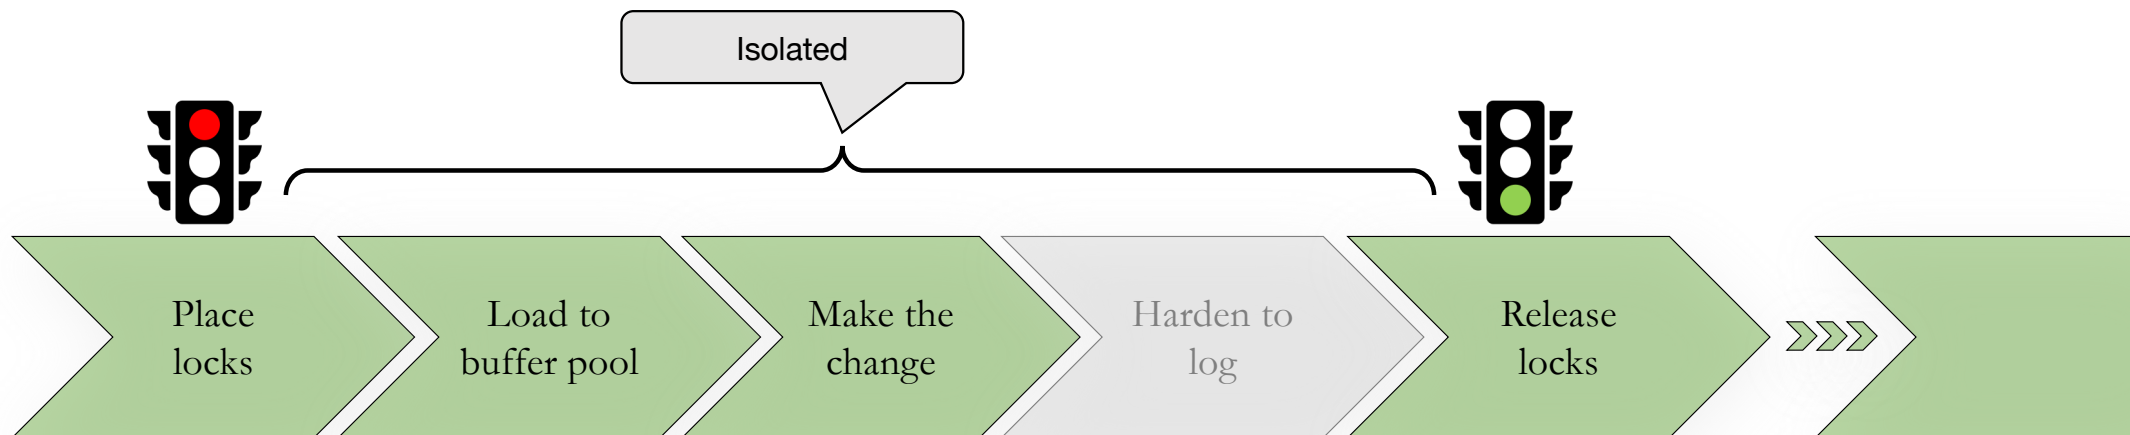
Durability



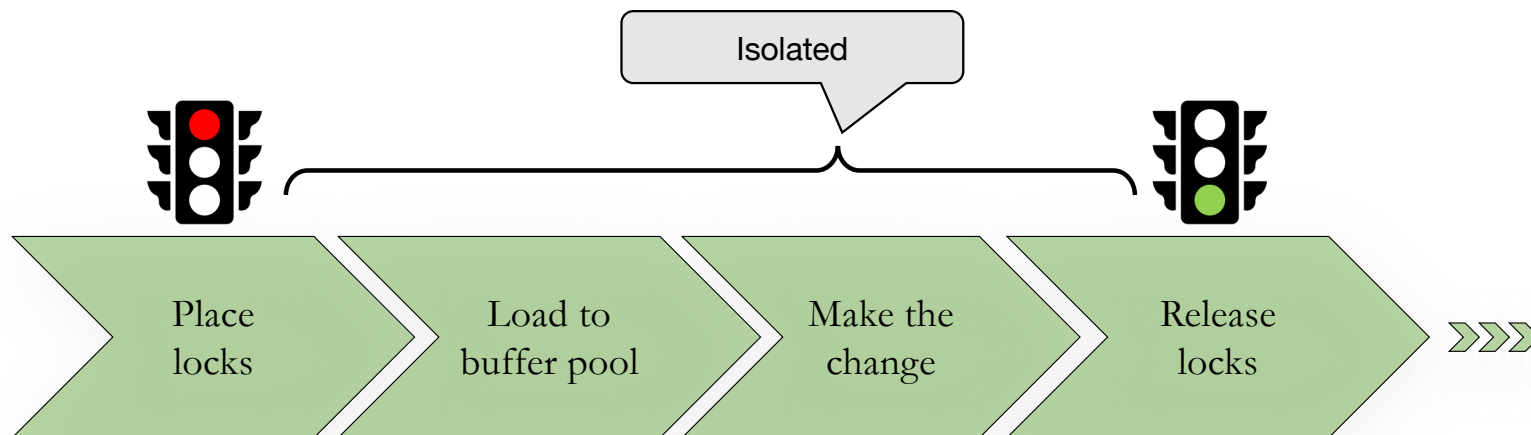
Durability



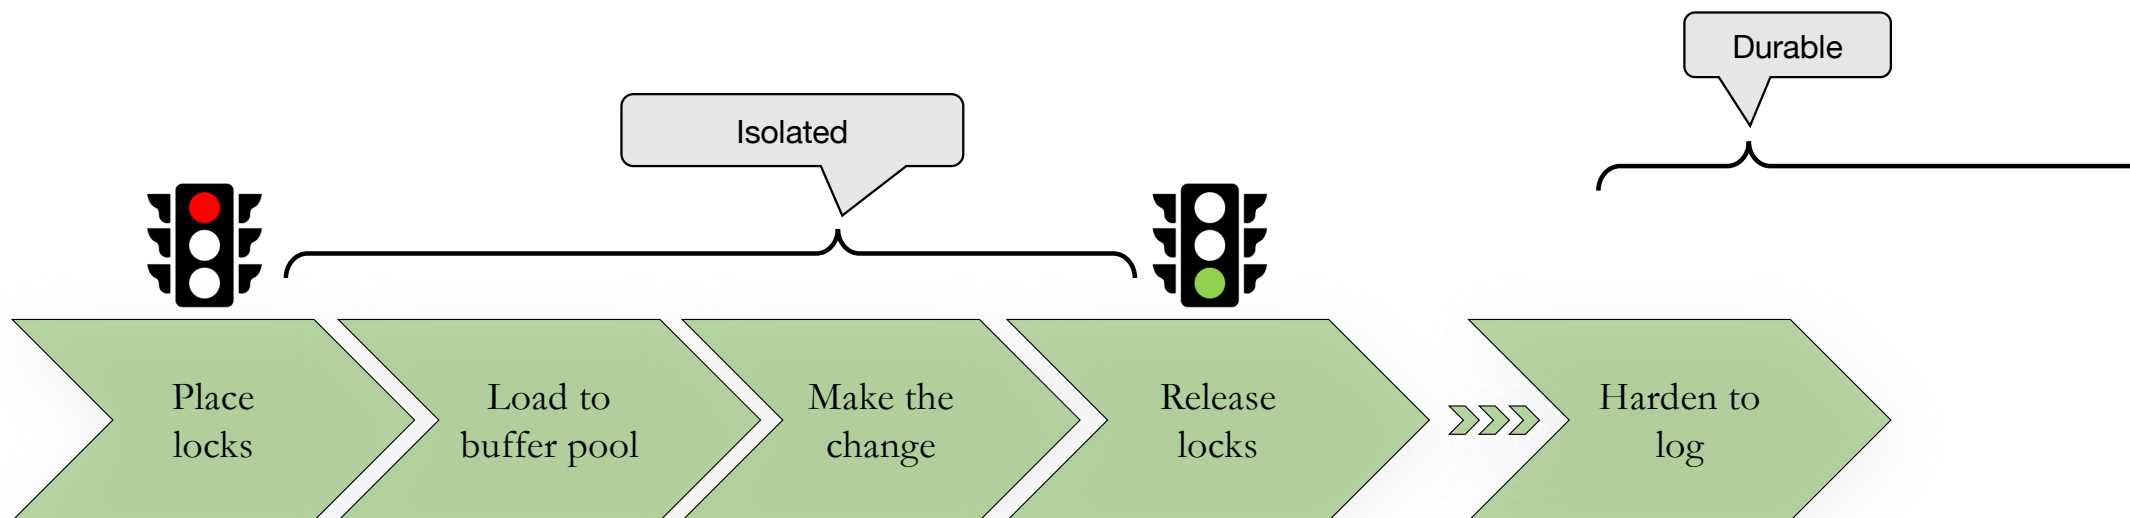
Delayed durability



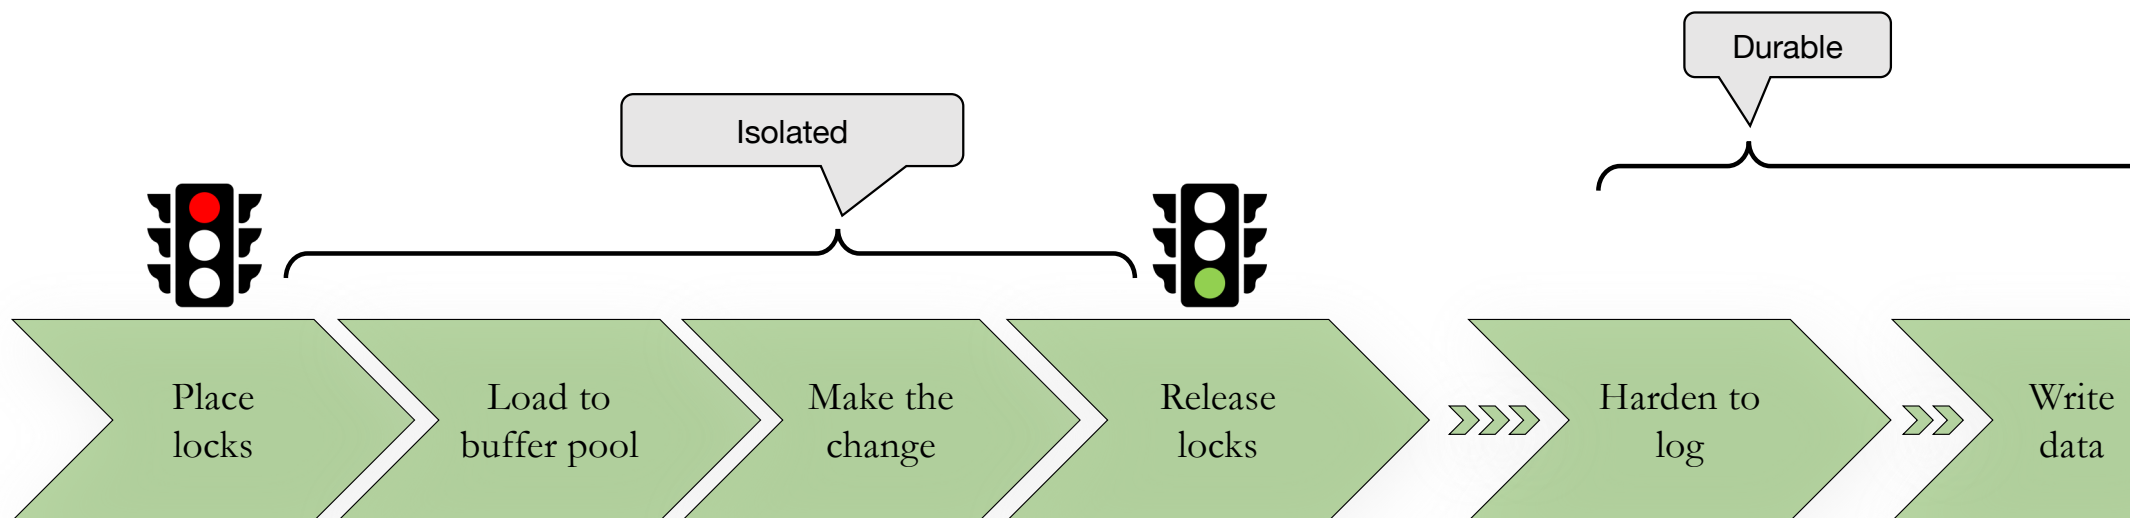
Delayed durability




Delayed durability



Delayed durability



Delayed durability



Isolated

Durable



Place
locks

Load to
buffer pool

Make the
change

Release
locks

Harden to
log

Write
data

Delayed durability

- Writes are batched into the transaction log
- Significantly reduced latency for tiny transactions
- Great if you don't mind losing the data – like a DW or a staging environment
- tempdb already uses a form of Delayed Durability under the hood



Takeaways and questions

- I was going to do this slide later, but here we are.



Please give me feedback



Oh, and the slides and scripts: github.com/sqlsunday/presentations

